

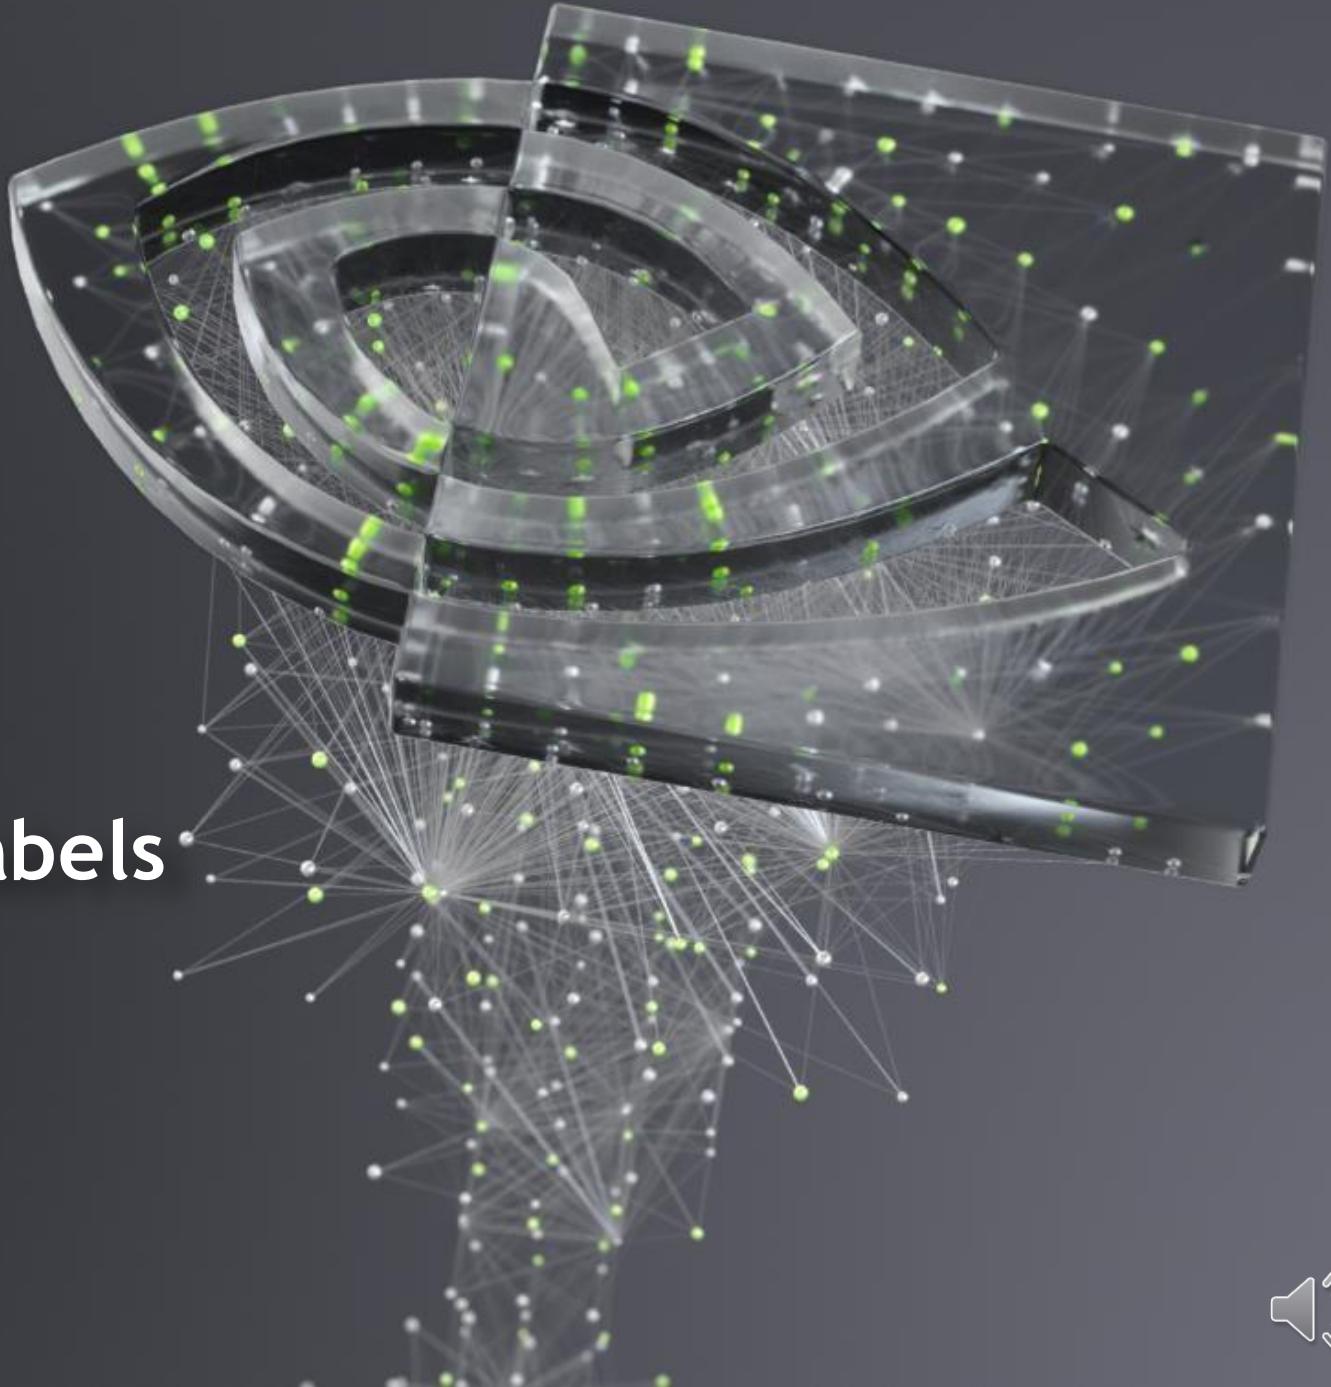


NVIDIA®

Learning with Imperfect Labels and Visual Data

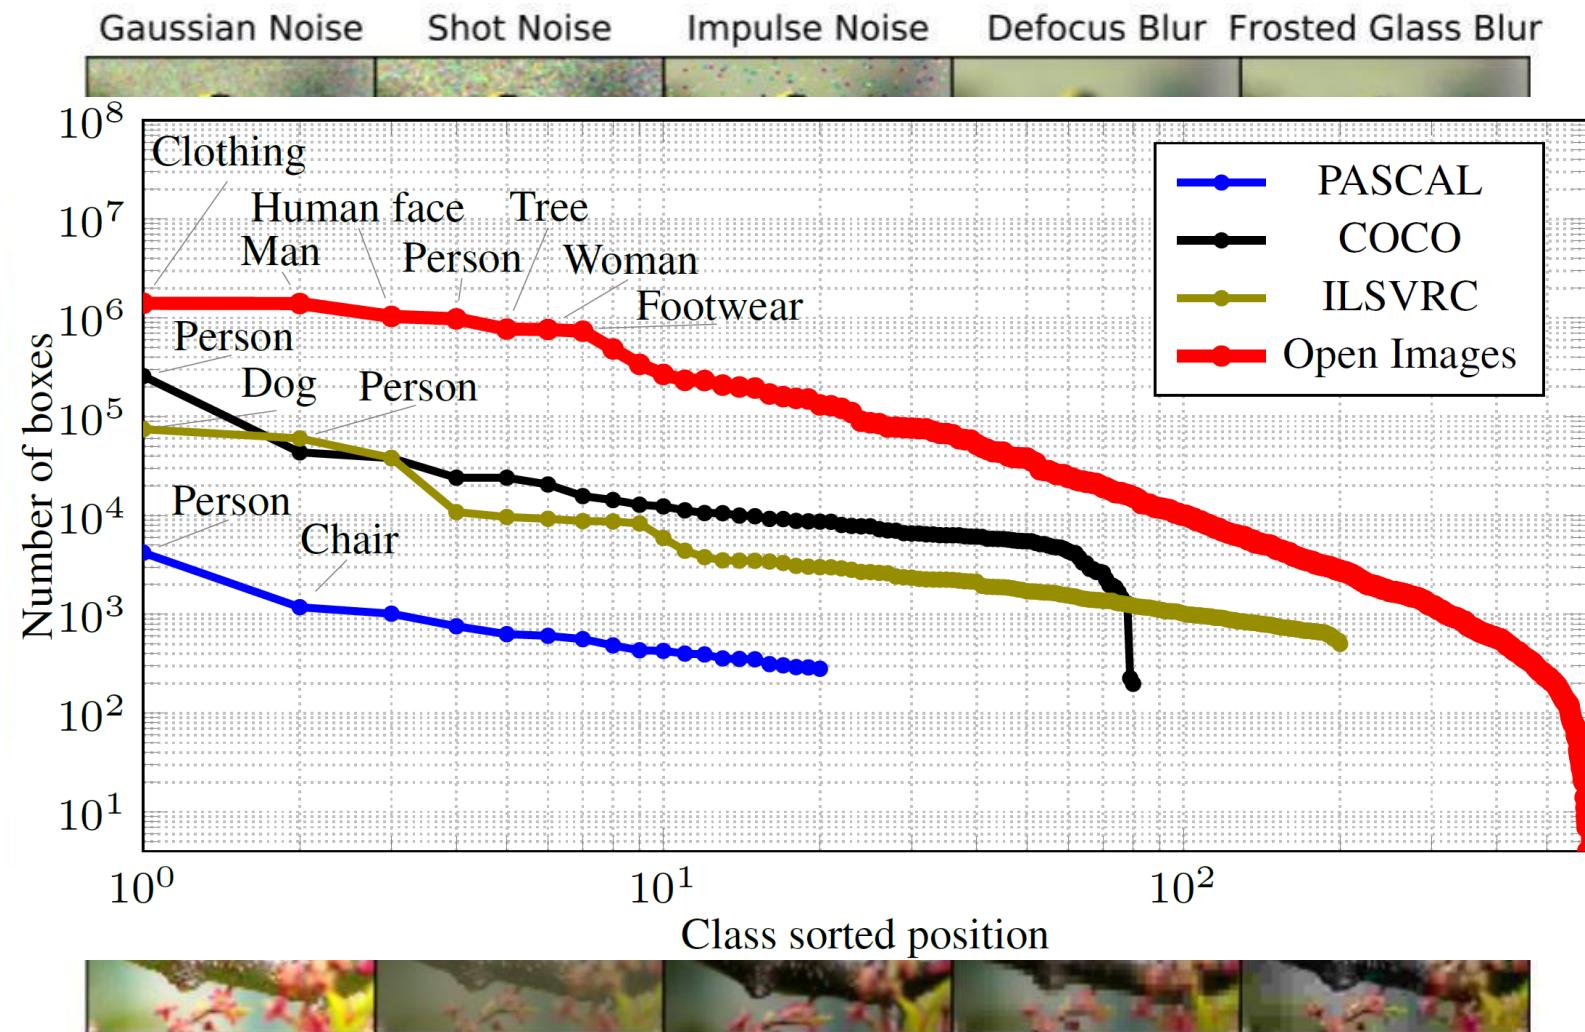
Zhidong Yu August 23rd 2020

New Frontiers for Learning with Limited Labels or Data



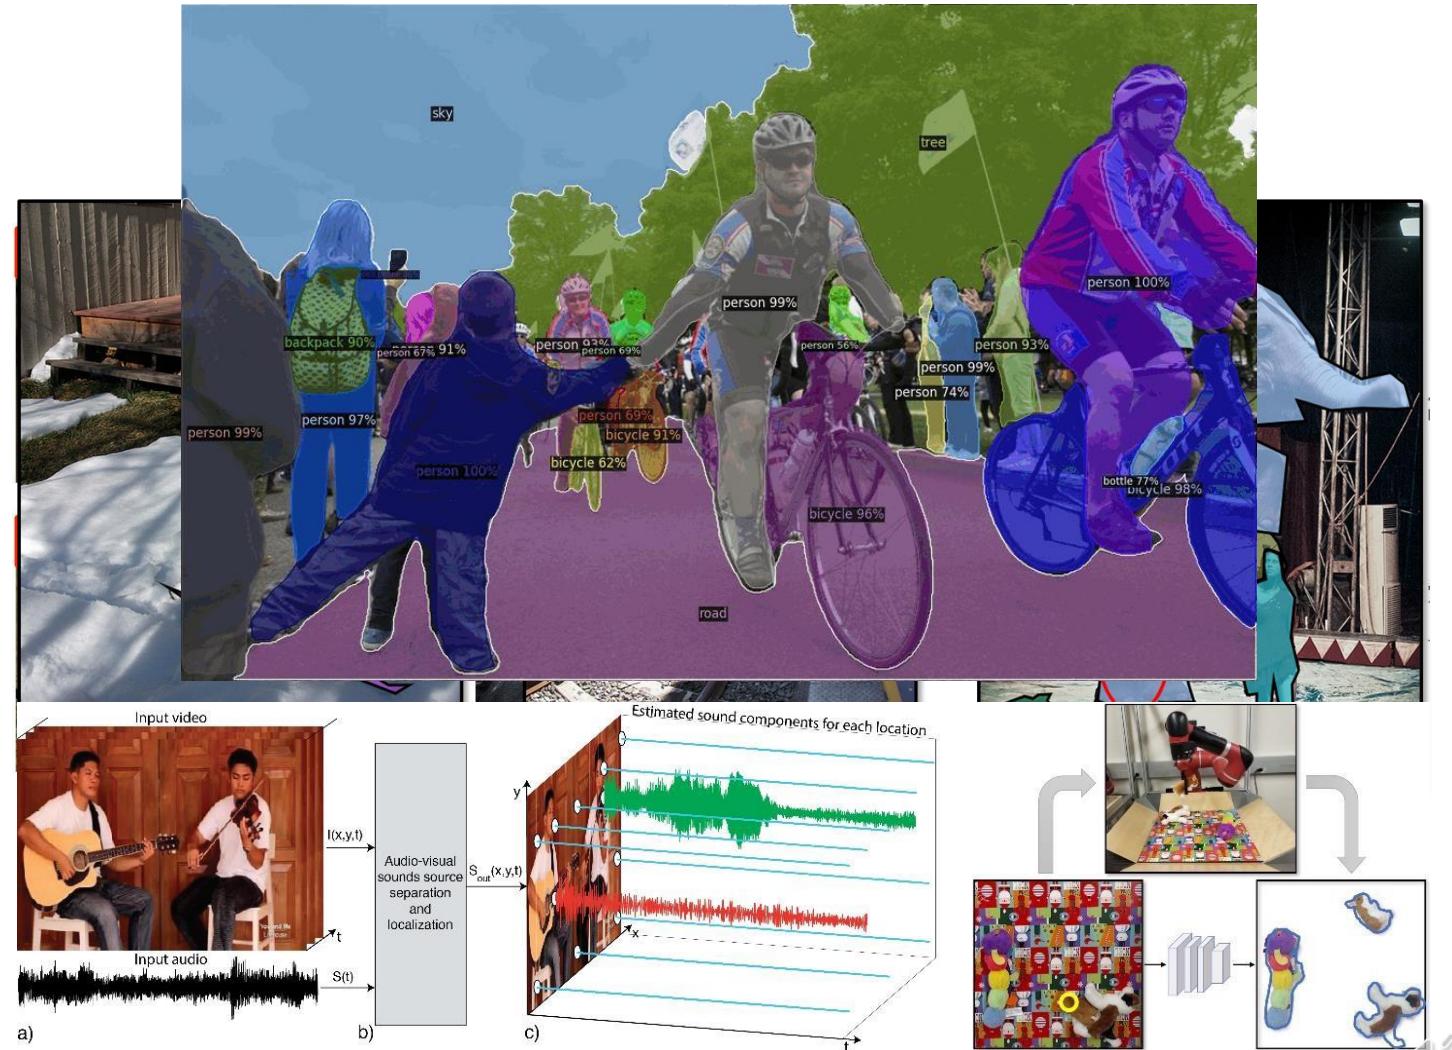
Challenge I - Real World Data Are Imperfect

- Domain gap
- Data bias
- Data noise
- Can be *Long tail*
- Can contain *Occlusions*
- Can be *Cluttered*
- Can be *Ambiguous*
- Can be *Multisensory*
- ...



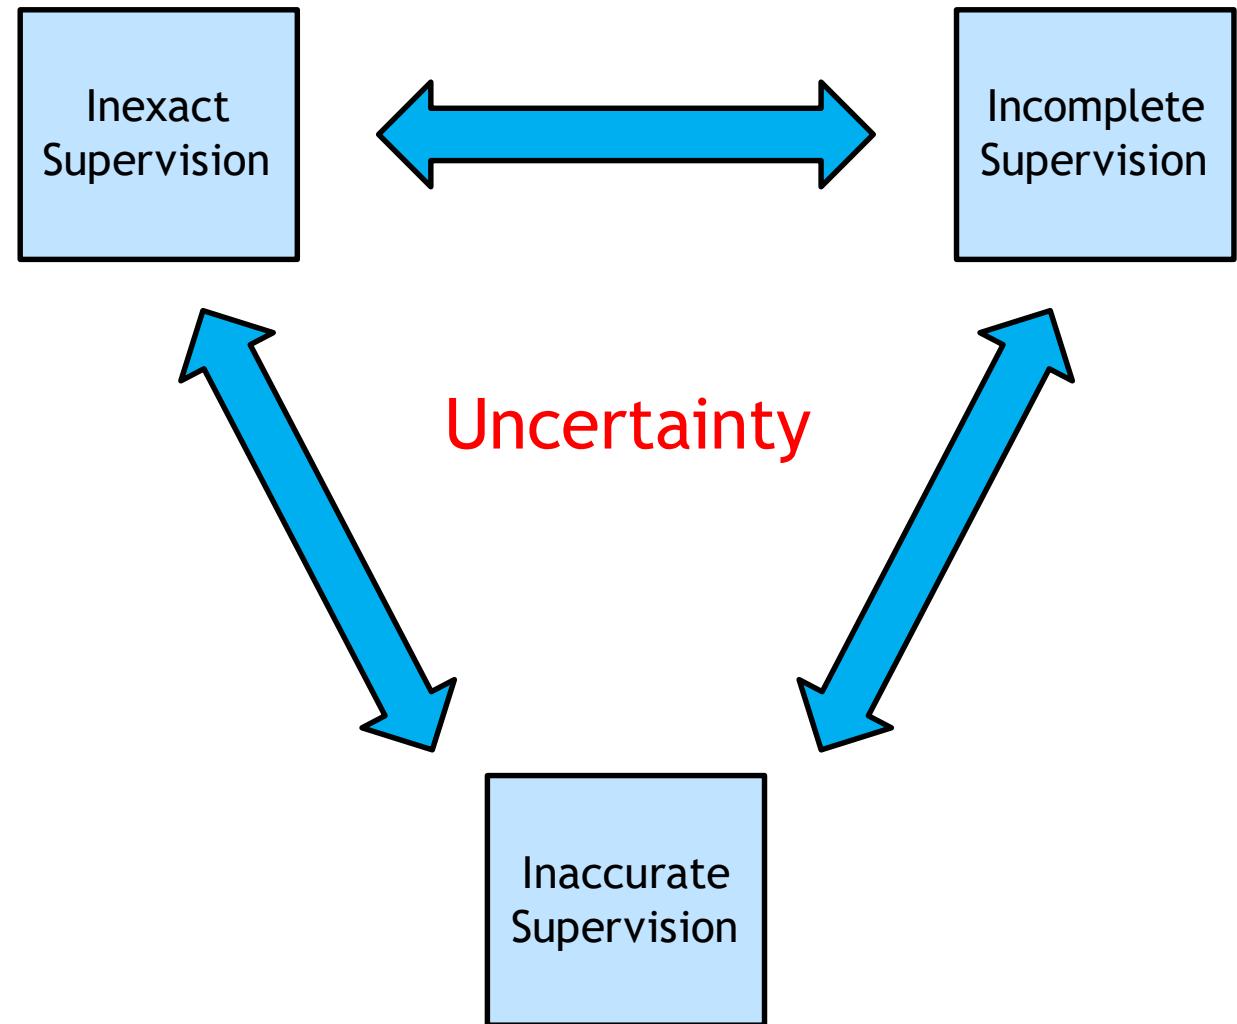
Challenge II - Real World Labels Are Imperfect

- Inexact (indirect) supervision
- Incomplete (limited) supervision
- Inaccurate (noisy) supervision
- Can be *Versatile*
- Can be *Multimodal*
- Can be *Sequential*
- Can be *Sparse*
- Can be *Interactive*
- ...

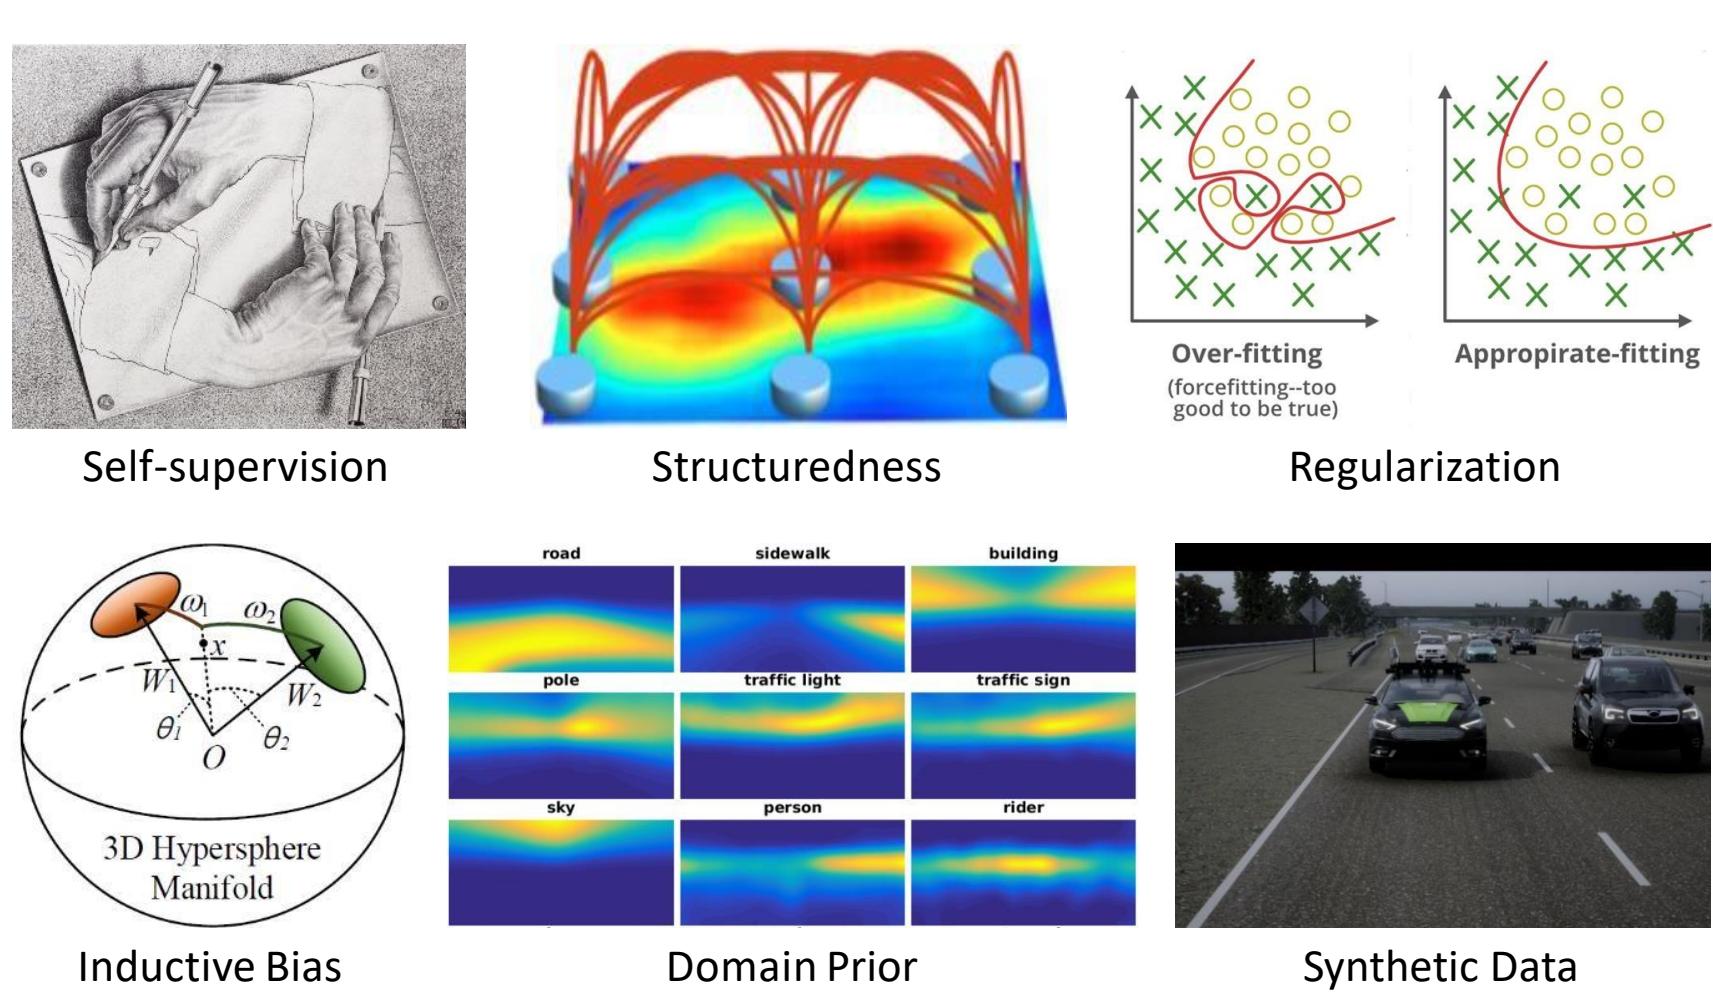
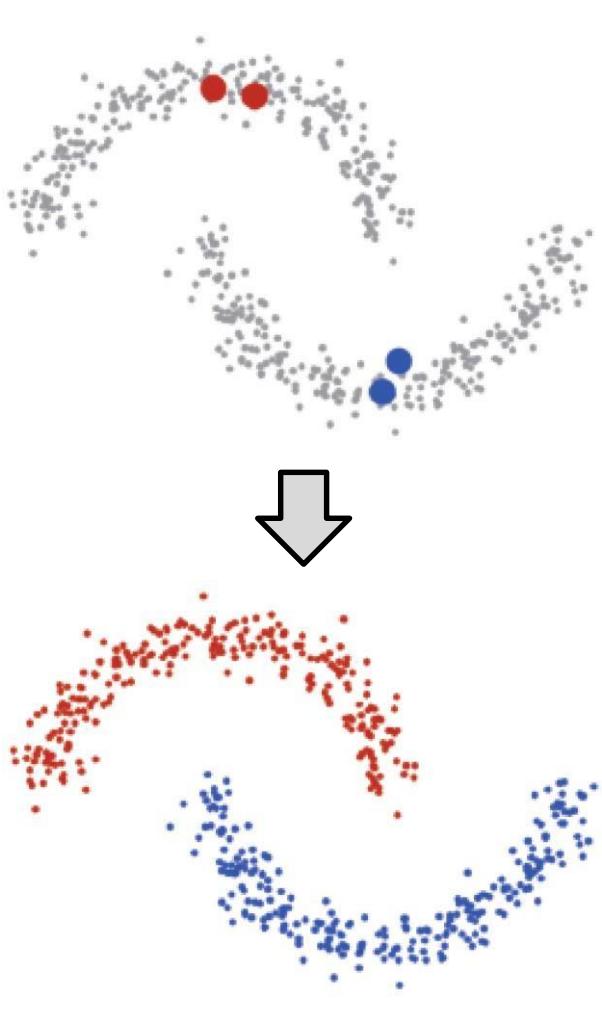


Goal and Challenge

- The ability to improve generalization by learning continuously on new data
- The ability to leverage diverse forms of weak supervision, or self-improving on unlabeled data
- The ability to overcome the **uncertainty** (ill-posed nature) due to the lack of constraints under partial supervision



Overcoming Uncertainties from Imperfect Labels





Part I: Learning on Unlabeled Data with Proxy-Label Approaches



Semi-Supervised Learning with Self-Ensembling

Labeled data $\mathcal{L} = \{(x_i, y_i)\}_{i=1}^L$, unlabeled data $\mathcal{U} = \{x_i\}_{i=L+1}^N$,
 $L \ll N$.

Learn $f : \mathcal{X} \rightarrow [0, 1]^K$ parameterized by $\theta \in \Theta$

$$\min_{\theta} \sum_{i=1}^L \ell(f(x_i; \theta), y_i) + \lambda R(\theta, \mathcal{L}, \mathcal{U}),$$

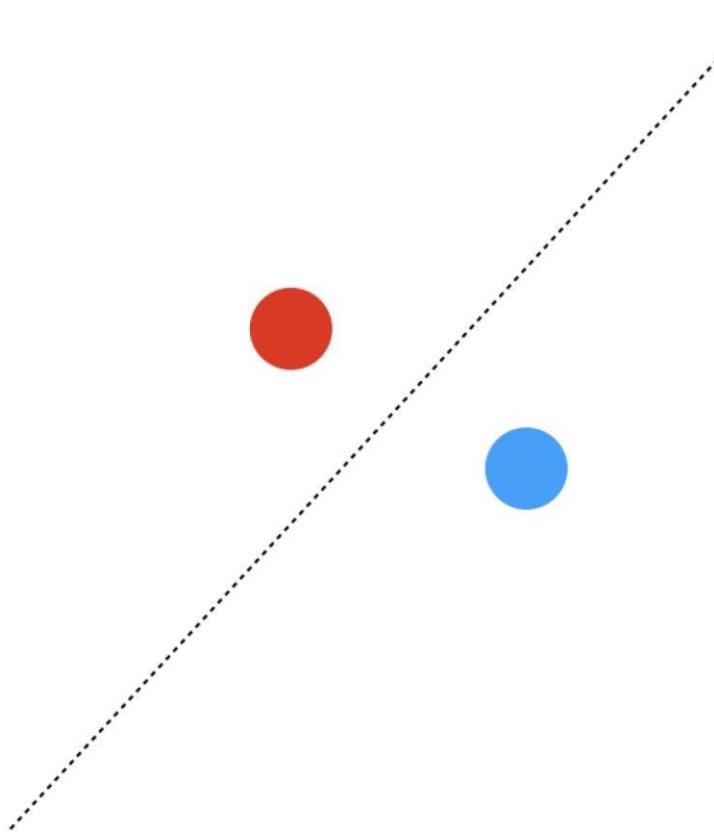
where R is the regularization term leveraging unlabeled data.

Core idea: Perturbing the input slightly and requiring that the prediction remains the same.

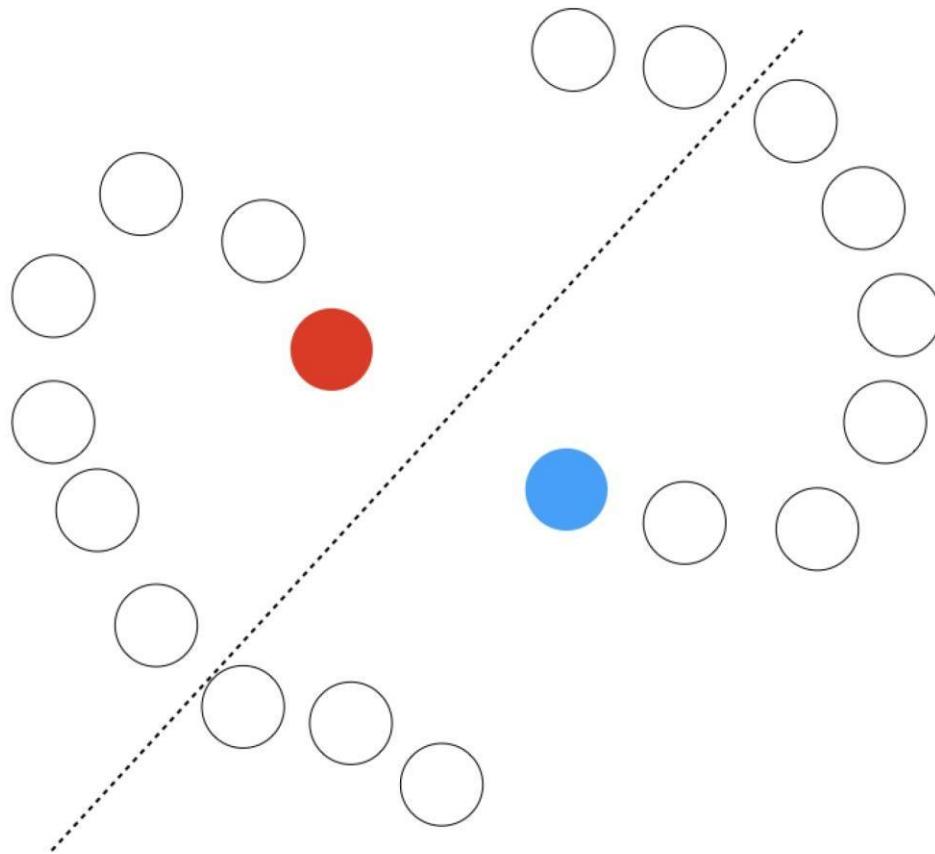
Define R as a consistency loss

$$R_C(\theta, \mathcal{L}, \mathcal{U}) = \sum_{i=1}^N \mathbb{E}_{\xi', \xi} d(\tilde{f}(x_i; \theta', \xi'), f(x_i; \theta, \xi))$$

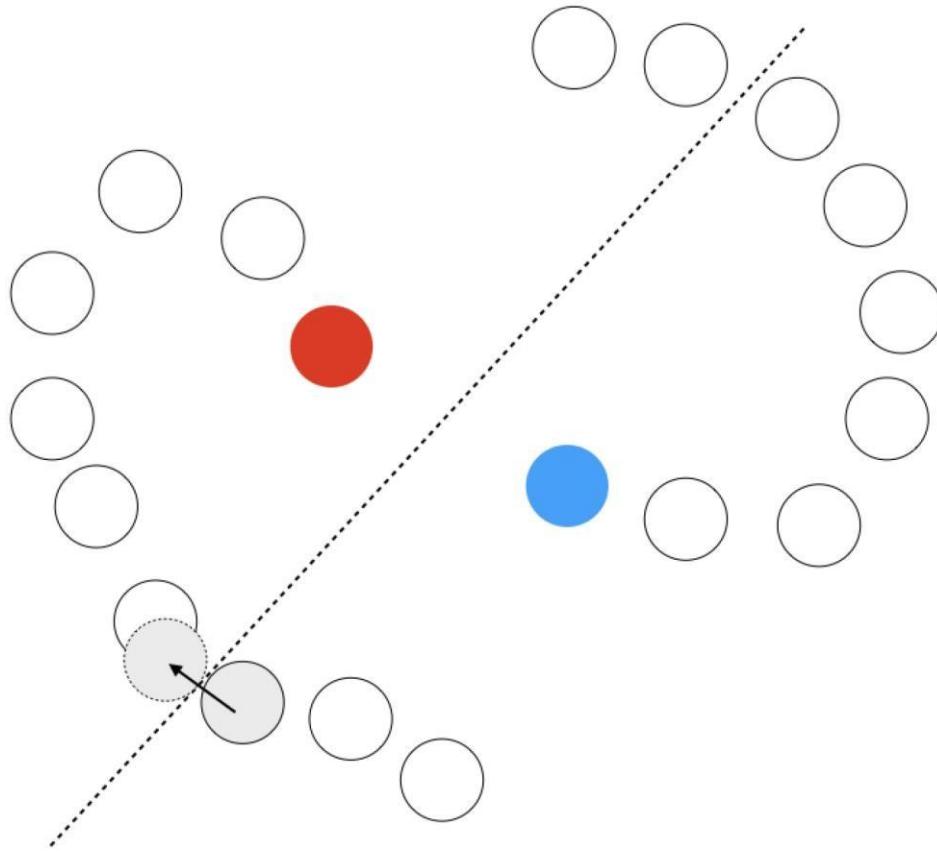
How Does Self-Ensembling Work?



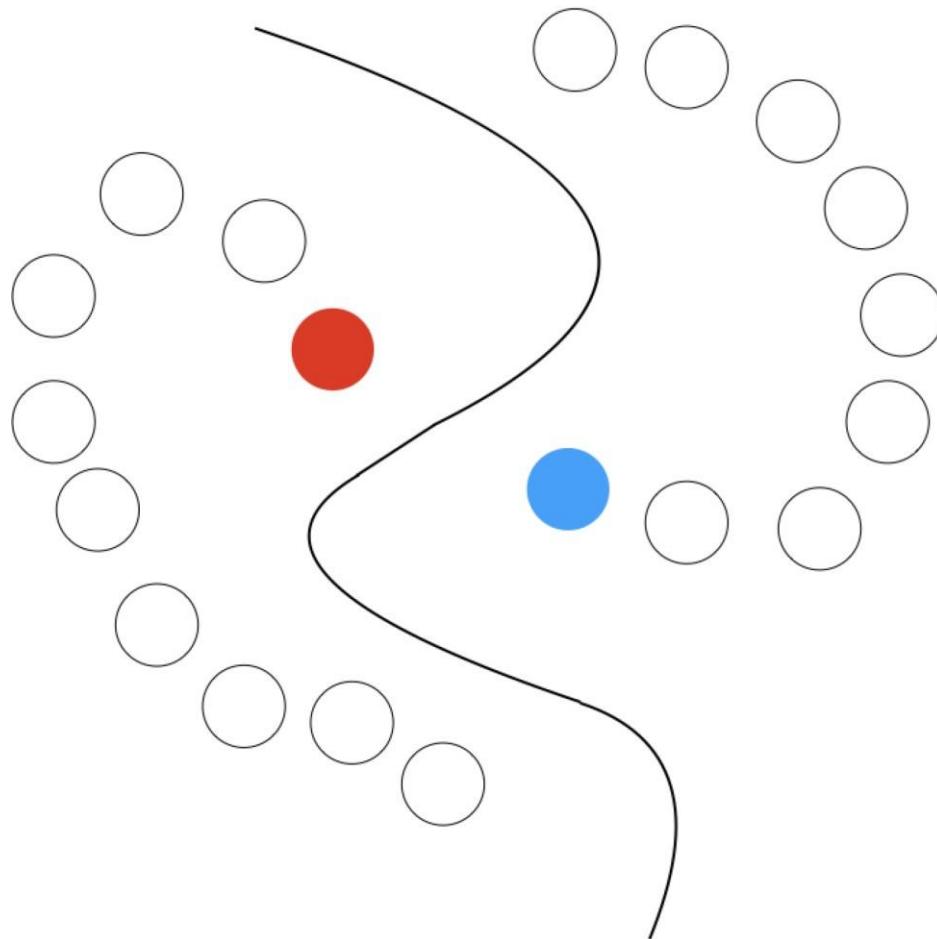
How Does Self-Ensembling Work?



How Does Self-Ensembling Work?



How Does Self-Ensembling Work?

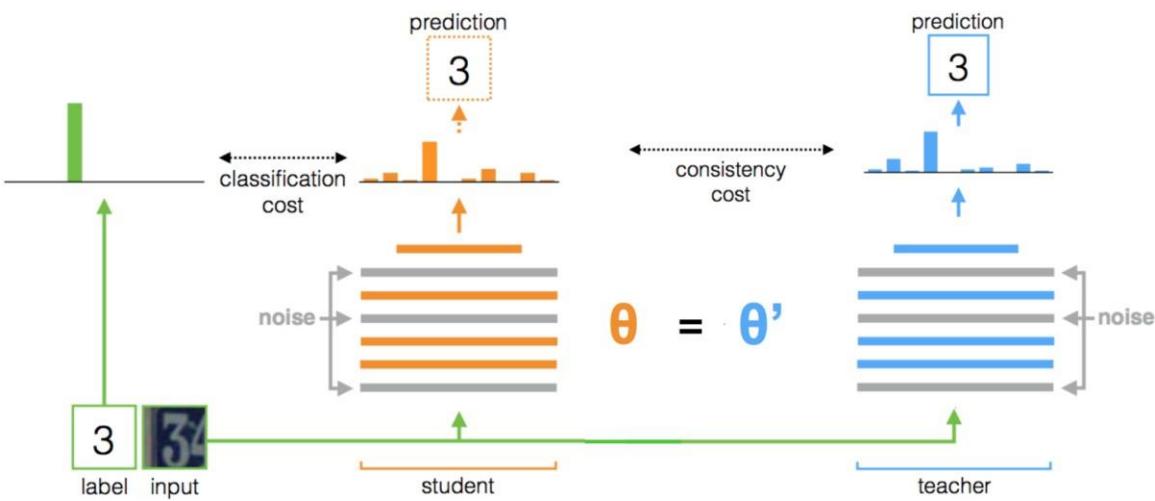


Self-Ensembling Methods

Π Model

Use a noisy teacher \tilde{f} to alleviate the bias in targets with $\theta' = \theta$

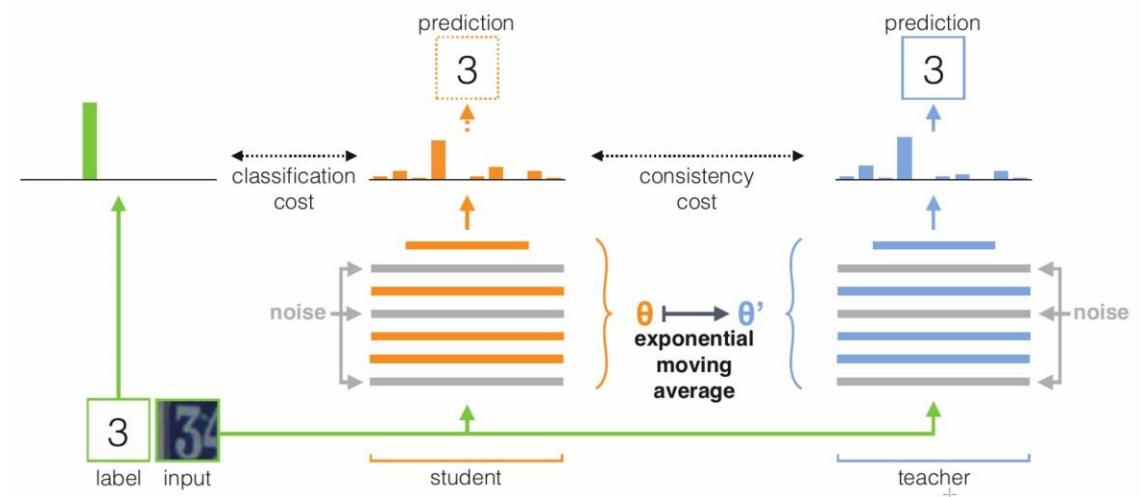
$$R_C(\theta, \mathcal{L}, \mathcal{U}) = \sum_{i=1}^N \mathbb{E}_{\xi', \xi} \|\tilde{f}(x_i; \theta, \xi') - f(x_i; \theta, \xi)\|^2,$$



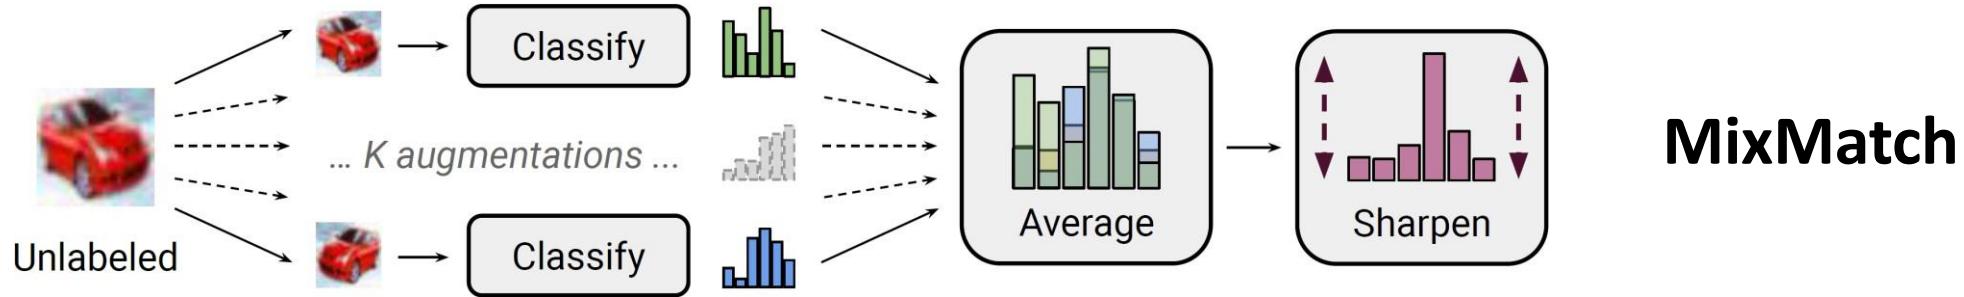
Mean Teacher Model

A better (more accurate) teacher using the EMA of parameters θ

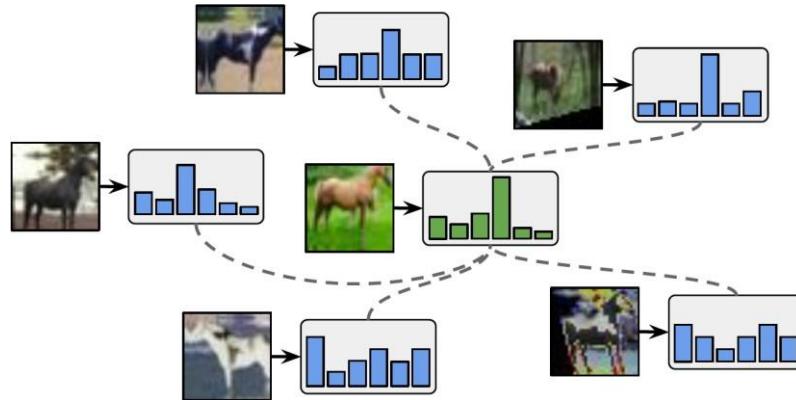
$$\theta'_t \leftarrow \alpha\theta'_{t-1} + (1 - \alpha)\theta_t$$



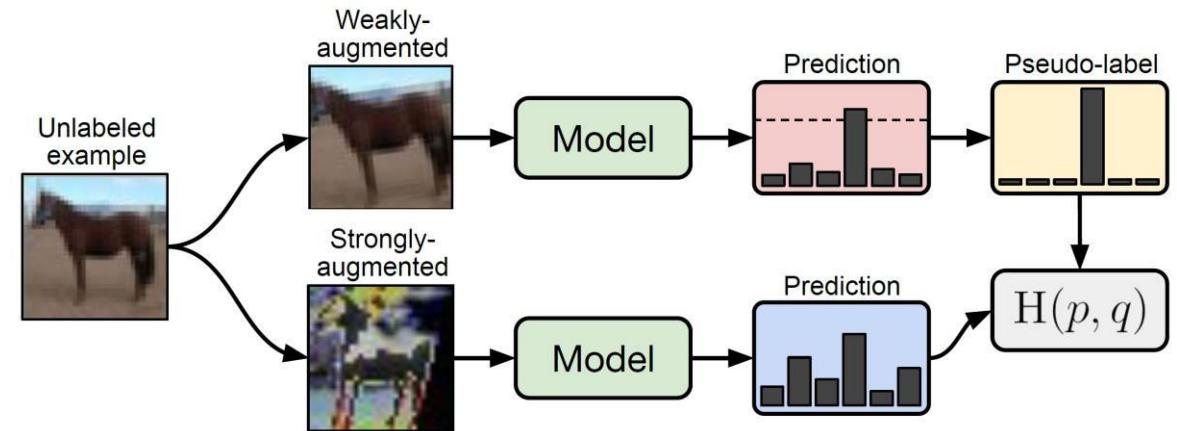
Semi-Sup Learning with Pseudo-Labels



ReMixMatch



FixMatch



Berthelot et al., MixMatch: A Holistic Approach to Semi-supervised Learning, NeurIPS19

Berthelot et al., ReMixMatch: Semi-supervised Learning with Distribution Alignment and Augmentation Anchoring, ICLR20

Sohn et al., FixMatch: Simplifying Semi-supervised Learning with Consistency and Confidence, arXiv:2001.07685

Unsupervised Domain Adaptation (UDA)

Image classification



Car



Semantic segmentation

Source Domain (Labeled)

Adaptation



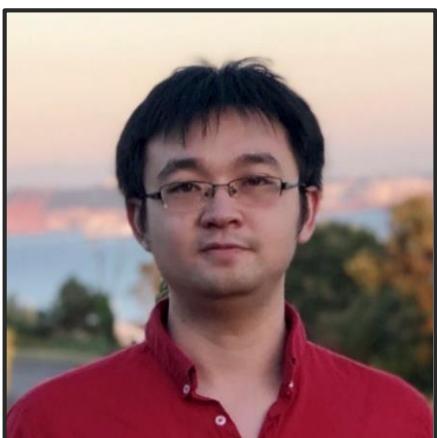
Adaptation



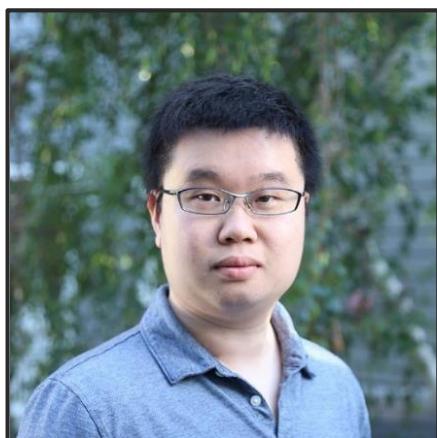
Target Domain (Unlabeled)

Domain Adaptation for Semantic Segmentation via Class-Balanced Self-Training (ECCV18)

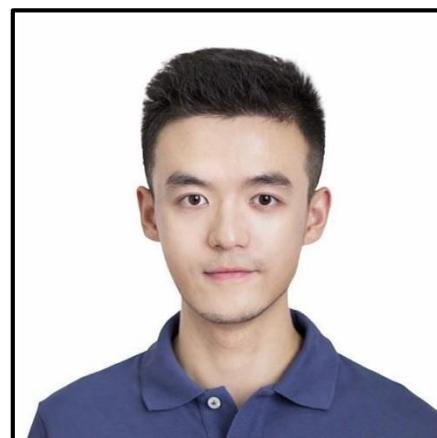
Confidence Regularized Self-Training (ICCV19)



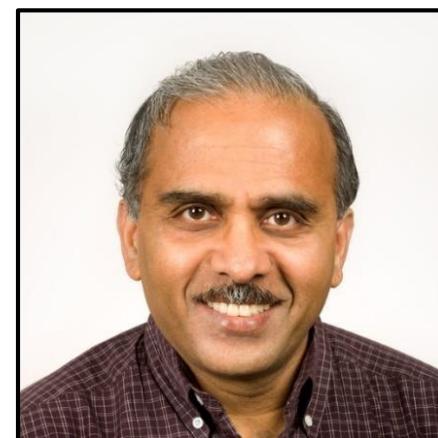
Yang Zou, CMU



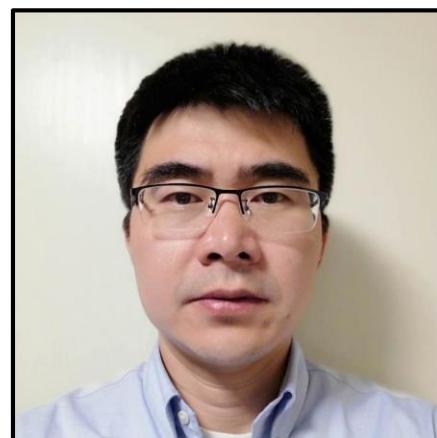
Zhidong Yu, NVIDIA



Xiaofeng Liu, CMU



Vijayakumar Bhagavatula,
CMU



Jinsong Wang, GM

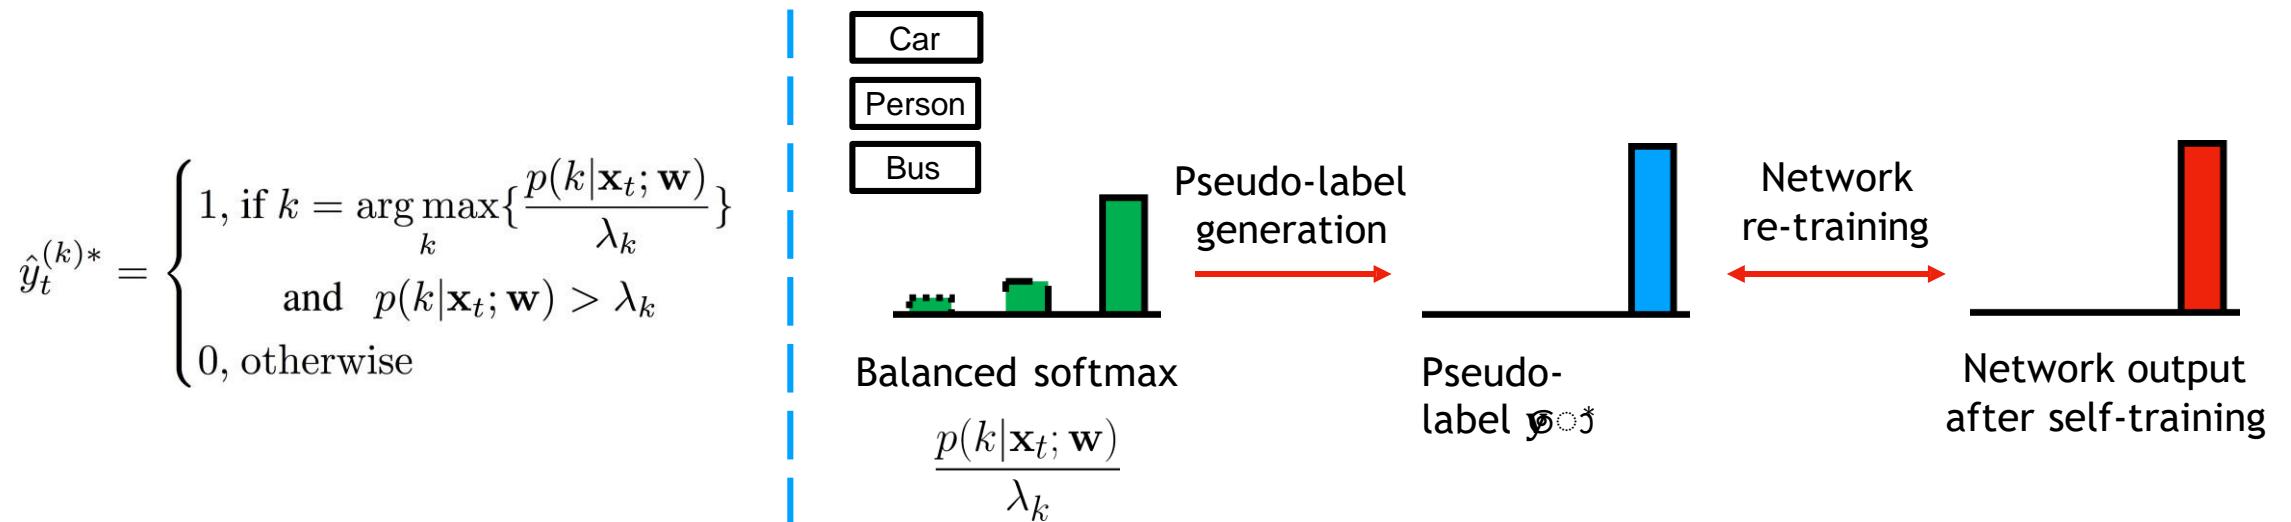
Class-Balanced Self-Training (CBST)

$$\min_{\mathbf{w}, \hat{\mathbf{Y}}_T} \mathcal{L}_{CB}(\mathbf{w}, \hat{\mathbf{Y}}_T) = - \sum_{s \in S} \sum_{k=1}^K y_s^{(k)} \log p(k|\mathbf{x}_s; \mathbf{w}) - \sum_{t \in T} \sum_{k=1}^K \hat{y}_t^{(k)} \log \frac{p(k|\mathbf{x}_t; \mathbf{w})}{\lambda_k}$$

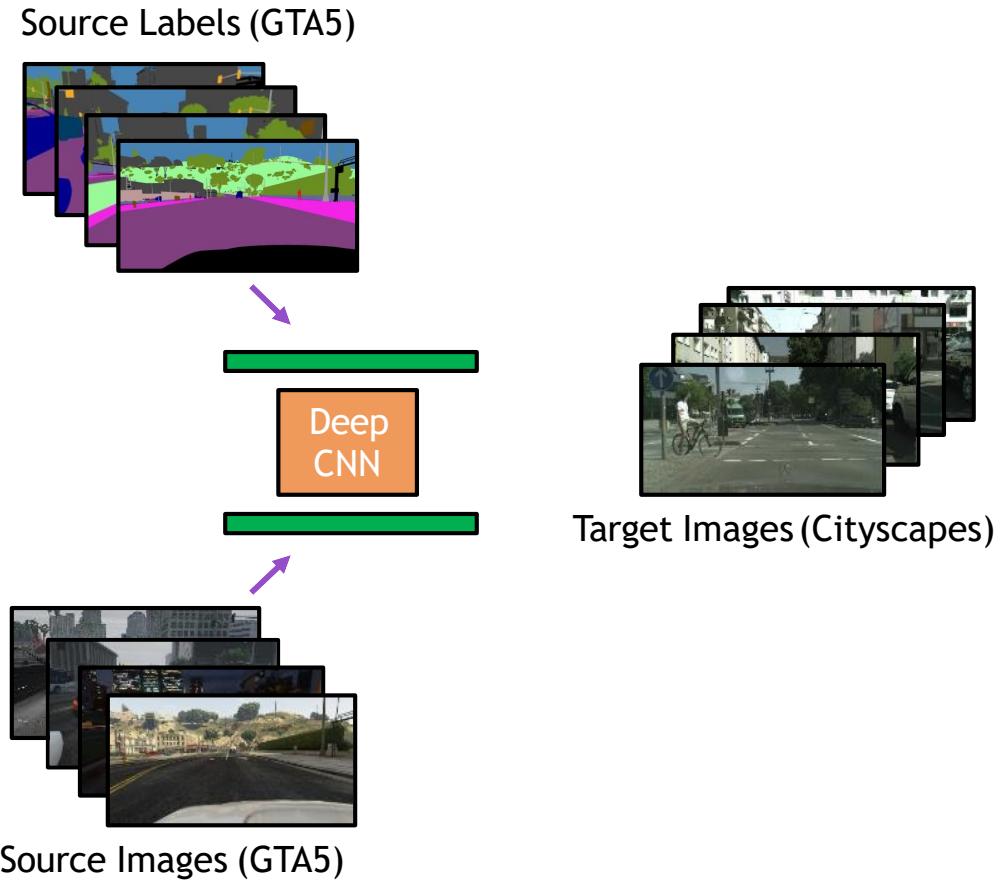
s.t. $\hat{\mathbf{y}}_t = (\hat{y}_t^{(1)}, \dots, \hat{y}_t^{(K)}) \in \Delta^{K-1} \cup \{\mathbf{0}\}, \forall t$

$\lambda_k > 0$

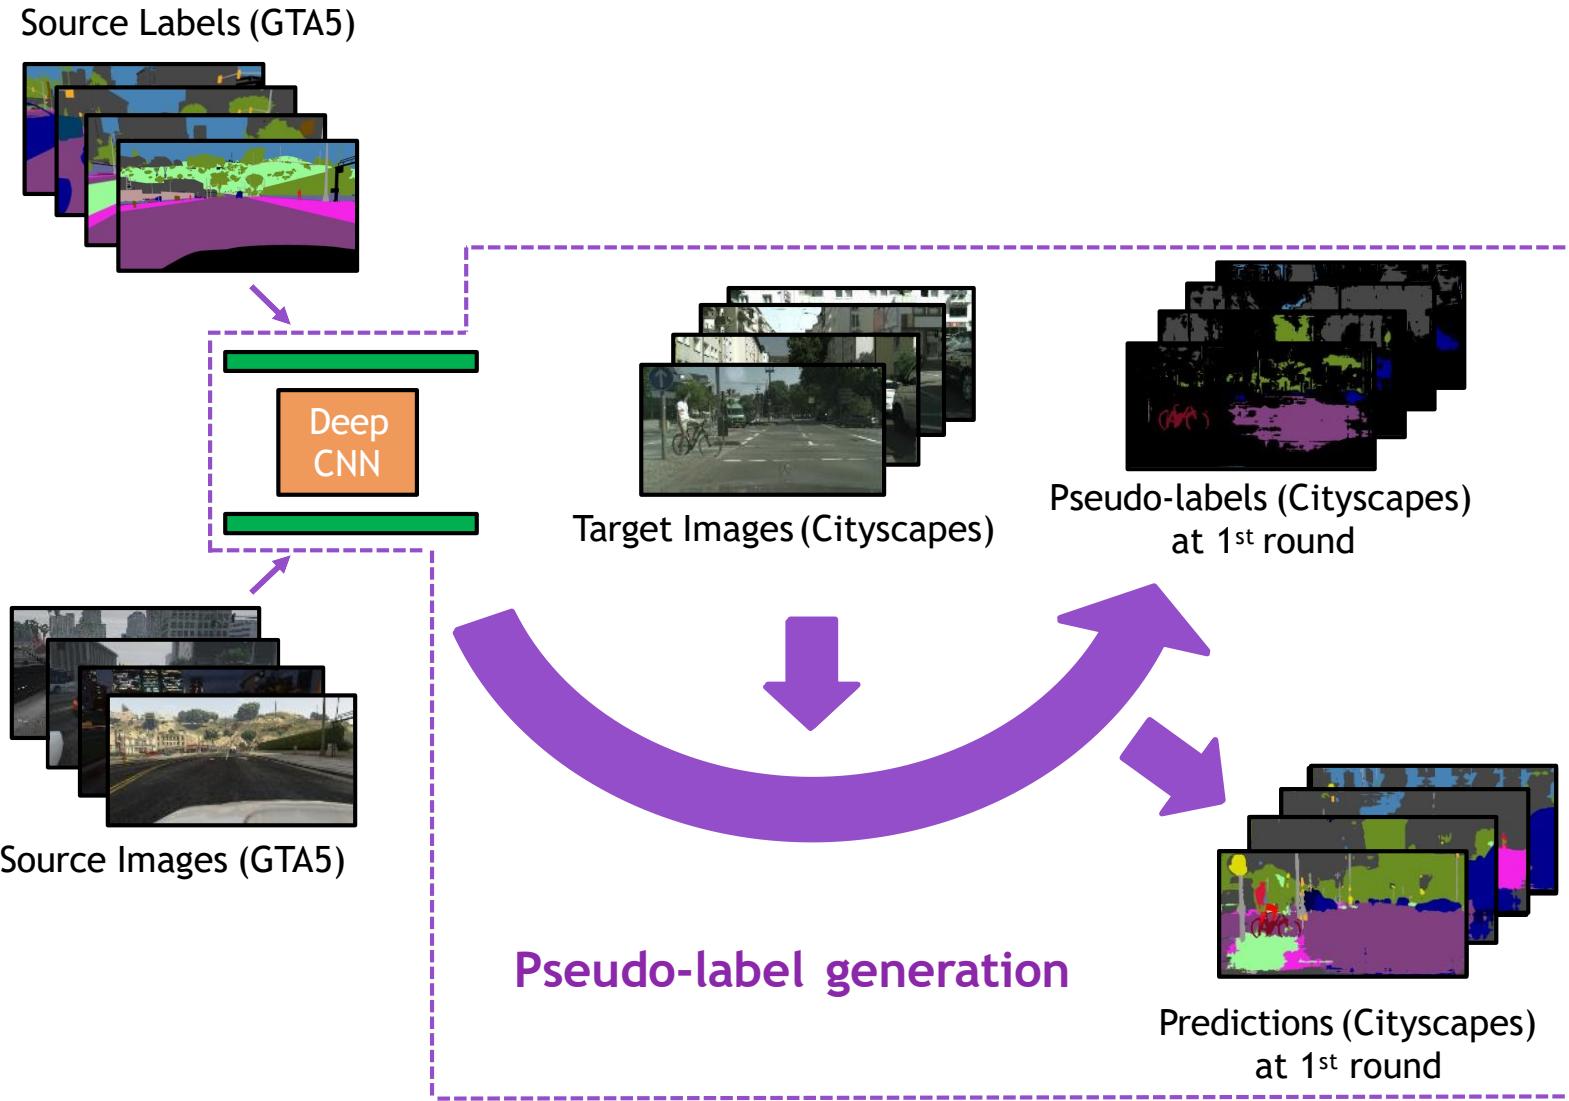
where: \mathbf{x} : input sample \mathbf{p} : class predication vector \mathbf{y} : label vector $\hat{\mathbf{y}}$: pseudo-label vector
 \mathbf{w} : network parameters s : source sample index t : target sample index Δ^{K-1} : probability simplex



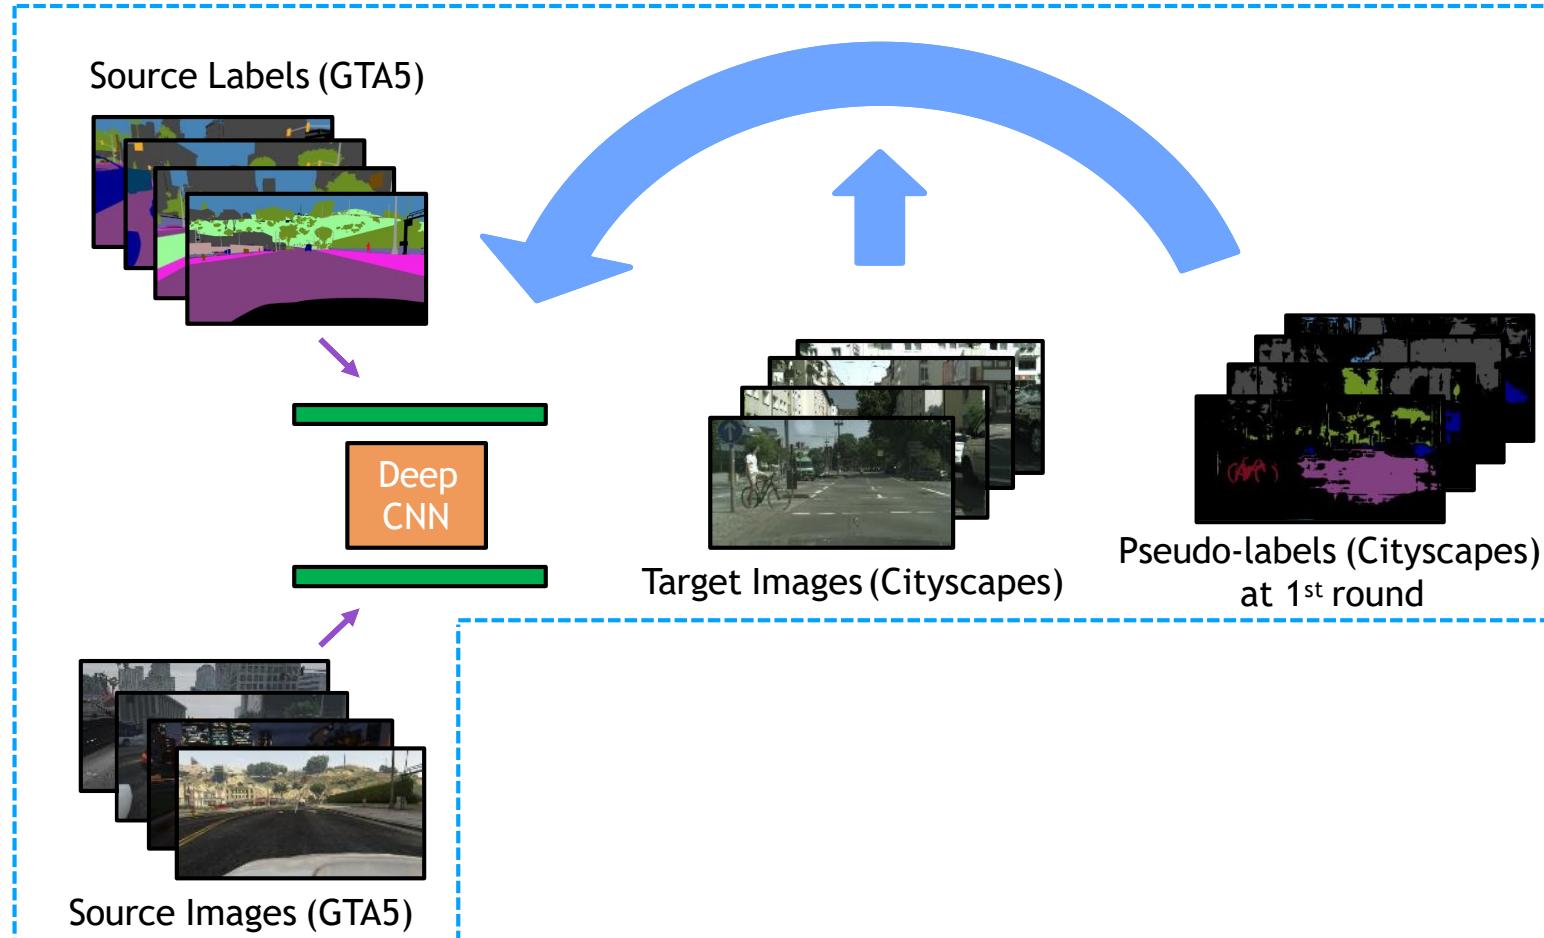
Domain Adaptation through Deep Self-Training



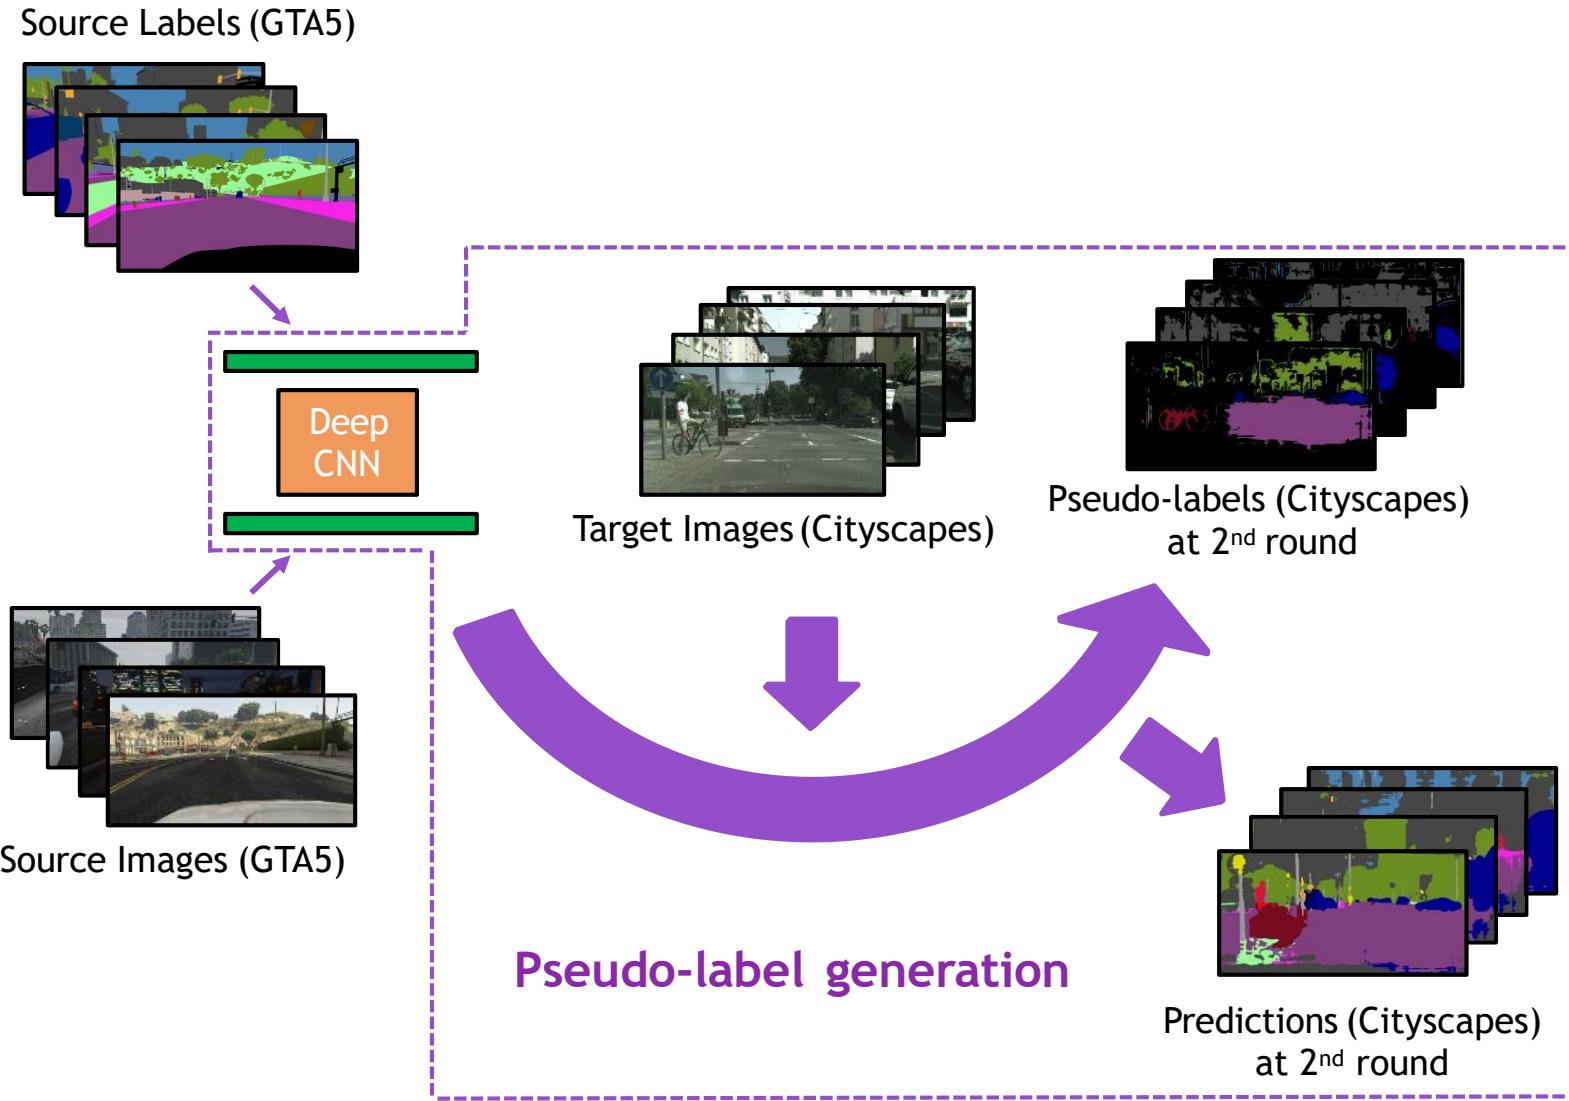
Domain Adaptation through Deep Self-Training



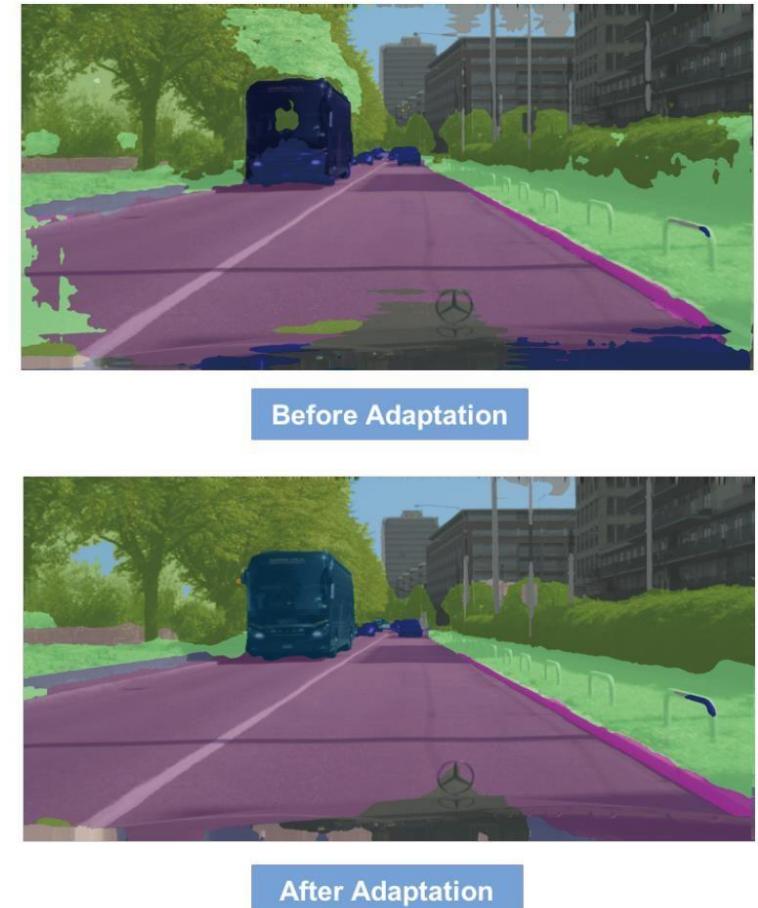
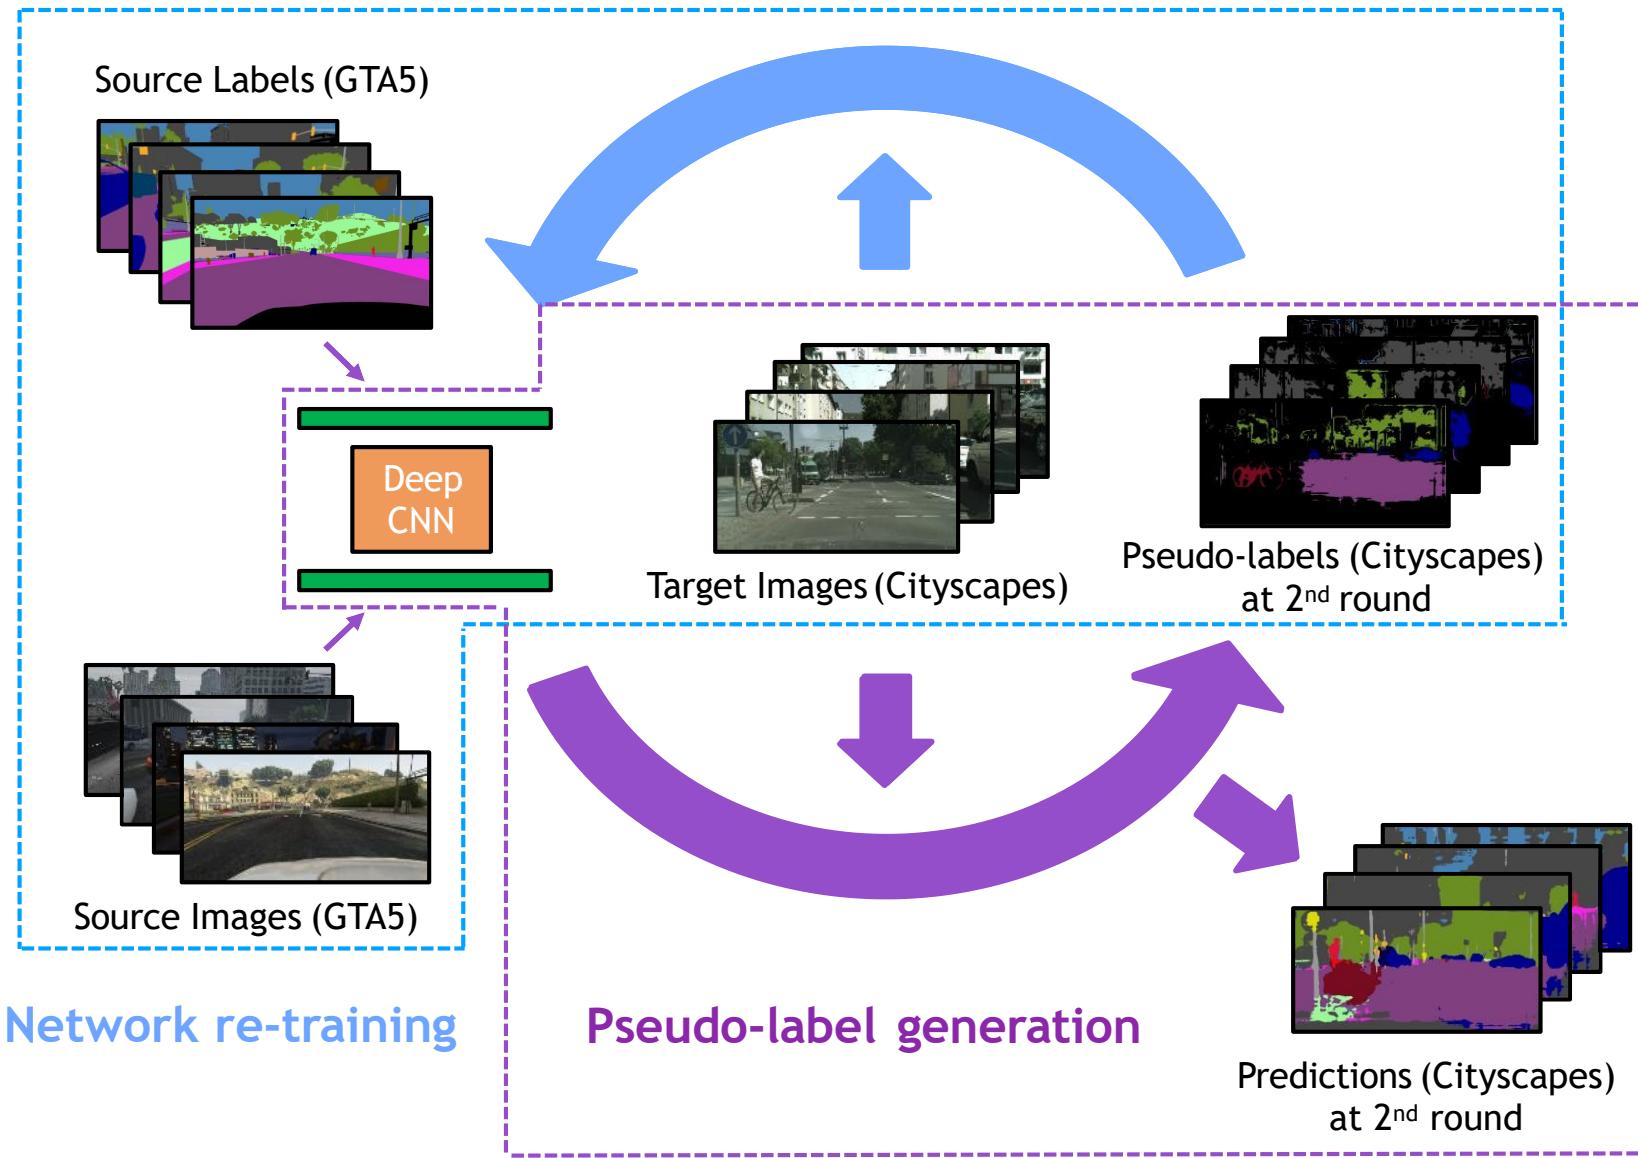
Domain Adaptation through Deep Self-Training



Domain Adaptation through Deep Self-Training



Domain Adaptation through Deep Self-Training



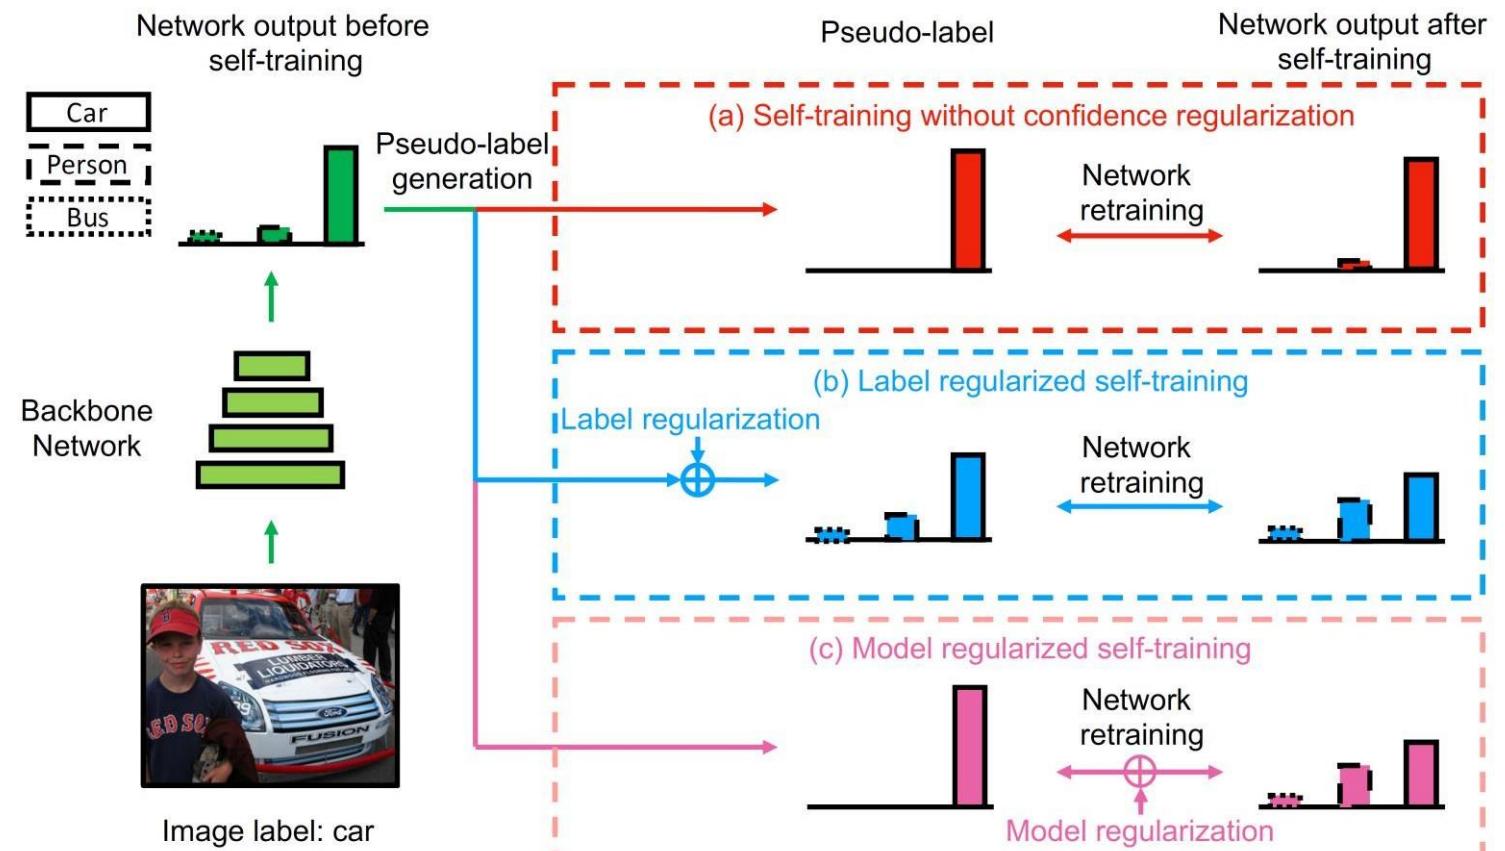
Confidence Regularized Self-Training (CRST)



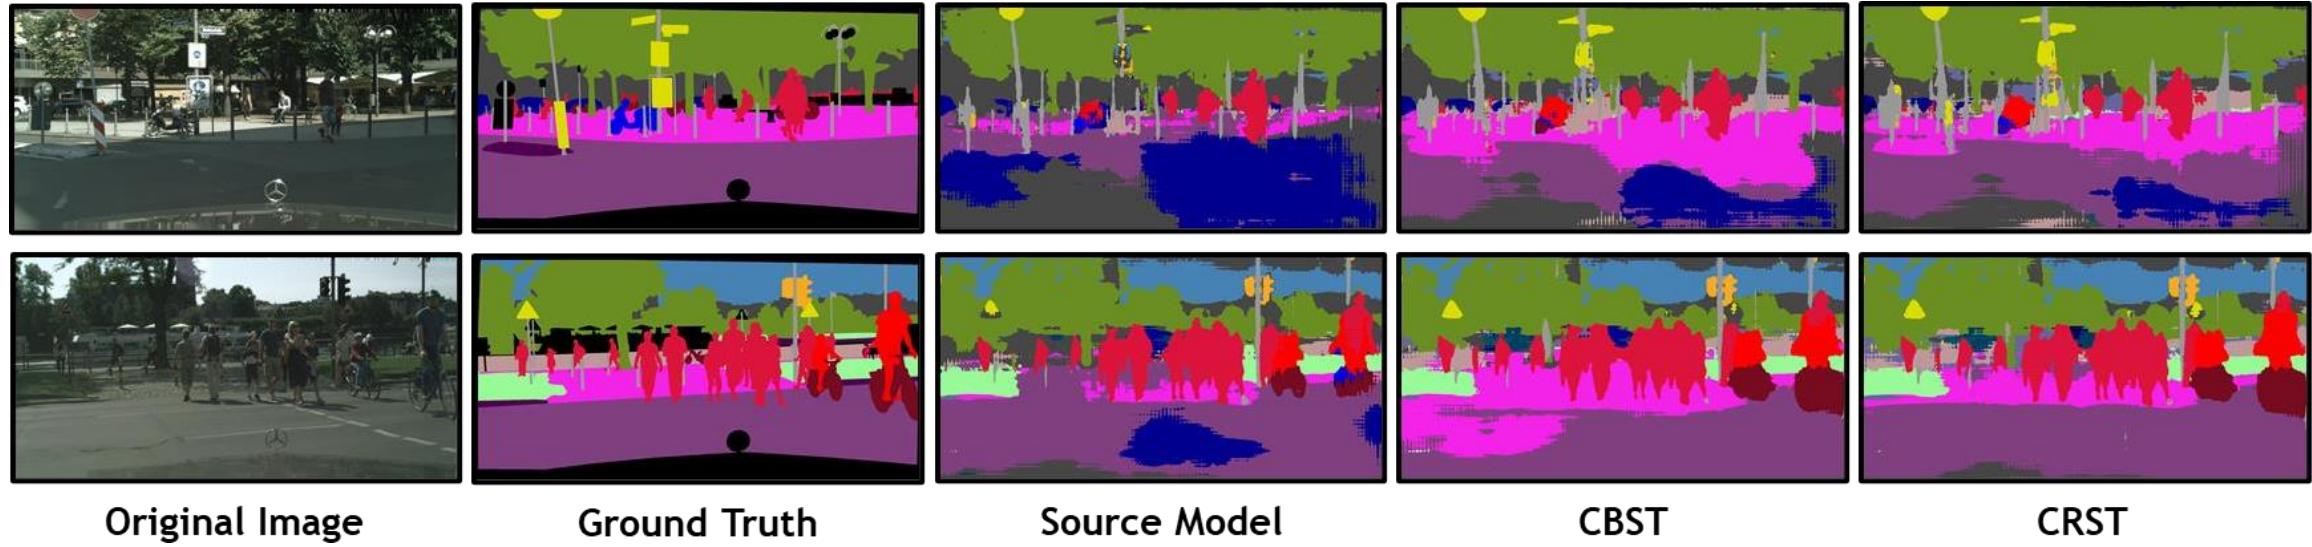
Sample from BDD100K
Green: Correct Red: Misclassified



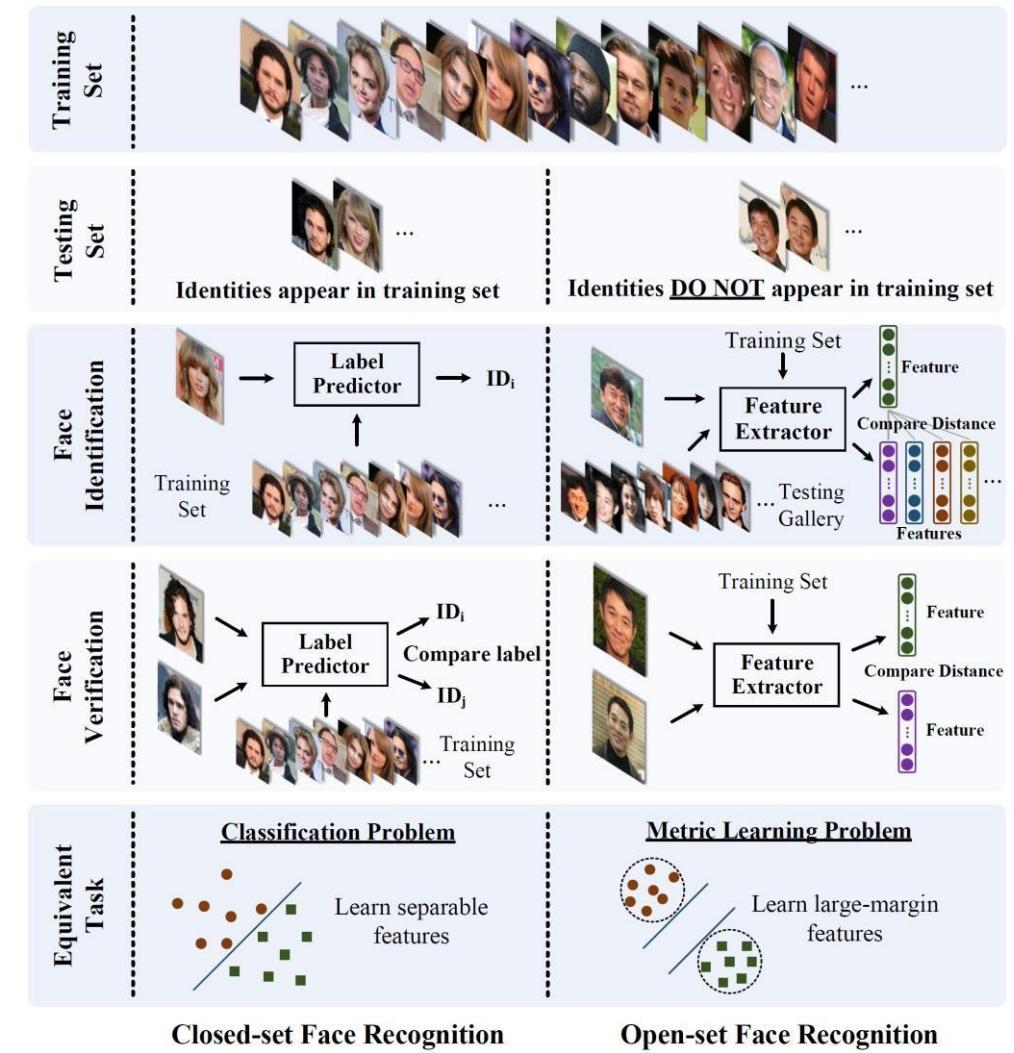
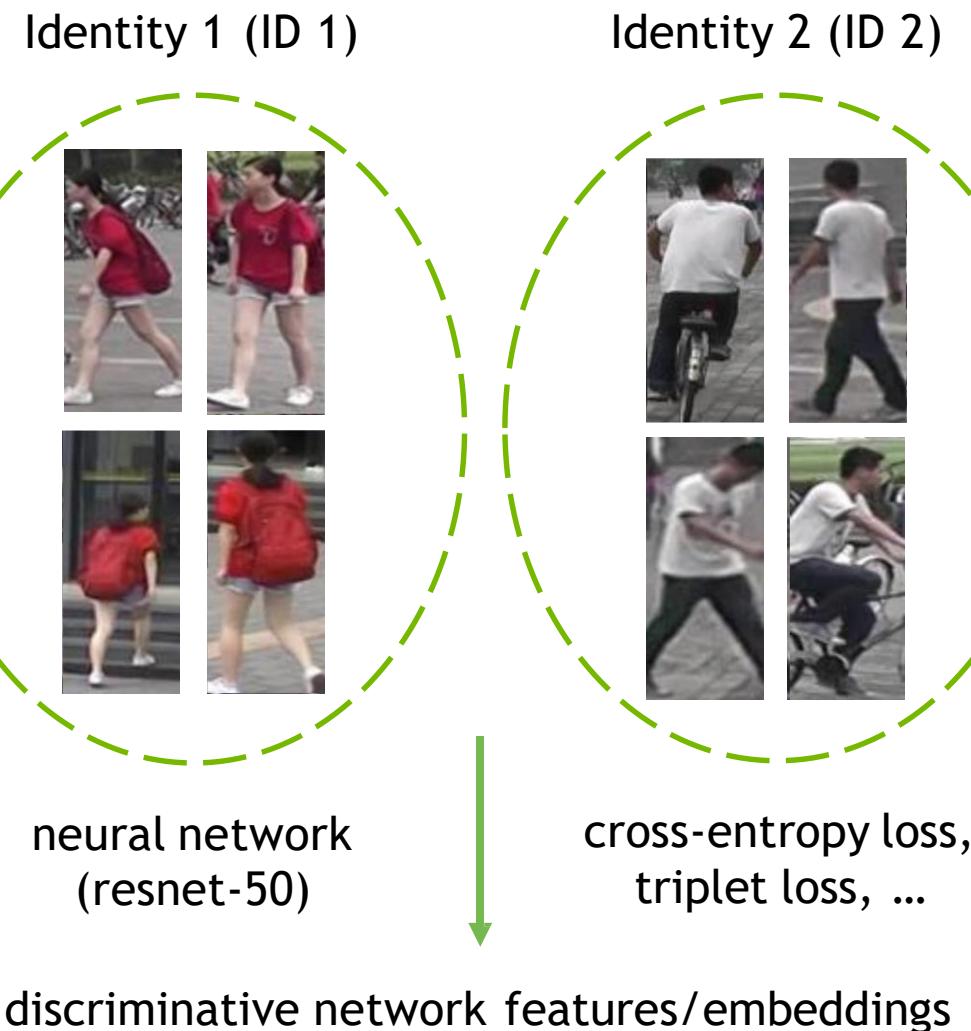
Samples from VisDA-17 (With label “Car”)



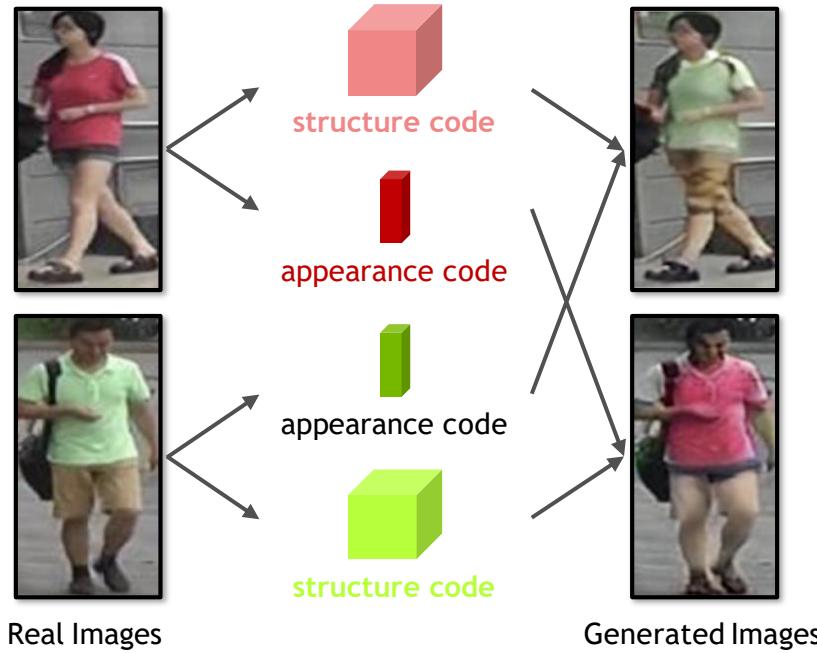
Qualitative Results: GTA5 -> Cityscapes



Person Re-ID: An Open Set Recognition Task



Cross-Domain Degeneration of DG-Net

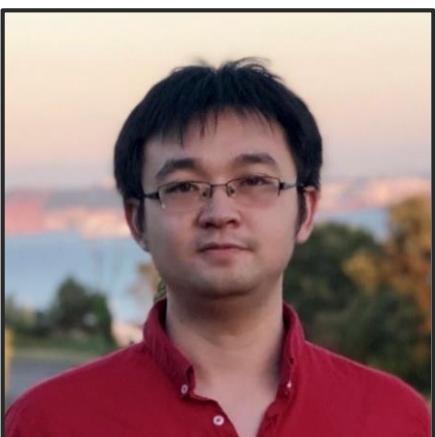


| Methods | Market-1501 | | DukeMTMC-reID | |
|--------------------|-------------|-------------|---------------|-------------|
| | Rank@1 | mAP | Rank@1 | mAP |
| Verif-Identif [55] | 79.5 | 59.9 | 68.9 | 49.3 |
| DCF [22] | 80.3 | 57.5 | - | - |
| SSM [2] | 82.2 | 68.8 | - | - |
| SVDNet [38] | 82.3 | 62.1 | 76.7 | 56.8 |
| PAN [57] | 82.8 | 63.4 | 71.6 | 51.5 |
| GLAD [47] | 89.9 | 73.9 | - | - |
| HA-CNN [24] | 91.2 | 75.7 | 80.5 | 63.8 |
| MLFN [4] | 90.0 | 74.3 | 81.0 | 62.8 |
| Part-aligned [37] | 91.7 | 79.6 | 84.4 | 69.3 |
| PCB [39] | 93.8 | 81.6 | 83.3 | 69.2 |
| Mancs [43] | 93.1 | 82.3 | 84.9 | 71.8 |
| DeformGAN [34] | 80.6 | 61.3 | - | - |
| LSRO [56] | 84.0 | 66.1 | 67.7 | 47.1 |
| Multi-pseudo [17] | 85.8 | 67.5 | 76.8 | 58.6 |
| PT [27] | 87.7 | 68.9 | 78.5 | 56.9 |
| PN-GAN [31] | 89.4 | 72.6 | 73.6 | 53.2 |
| FD-GAN [10] | 90.5 | 77.7 | 80.0 | 64.5 |
| Ours | 94.8 | 86.0 | 86.6 | 74.8 |

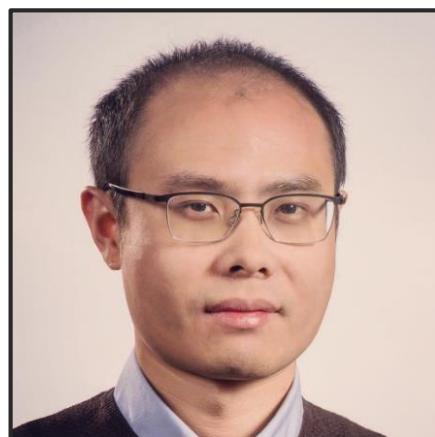
| Appearance Space | Structure Space |
|--|--|
| clothing/shoes color, texture and style, other id-related cues, etc. | body size, hair, carrying, pose, background, position, viewpoint, etc. |

| Metric | Market→Duke | Duke→Market | MSMT→Market | Market→MSMT | MSMT→Duke | Duke→MSMT |
|--------|-------------|-------------|-------------|-------------|-----------|-----------|
| Rank@1 | 42.6 | 56.1 | 61.8 | 17.1 | 61.9 | 20.6 |
| mAP | 24.3 | 26.8 | 33.6 | 5.4 | 40.7 | 6.4 |

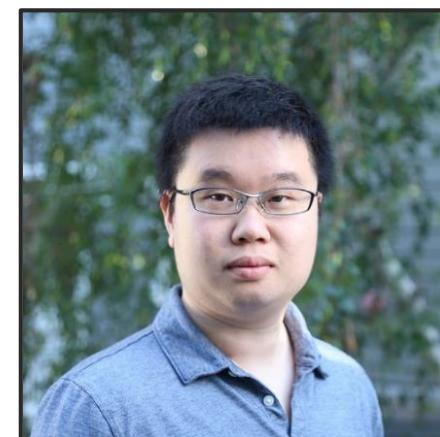
Joint Disentangling and Adaptation for Cross-Domain Person Re-Identification (ECCV20)



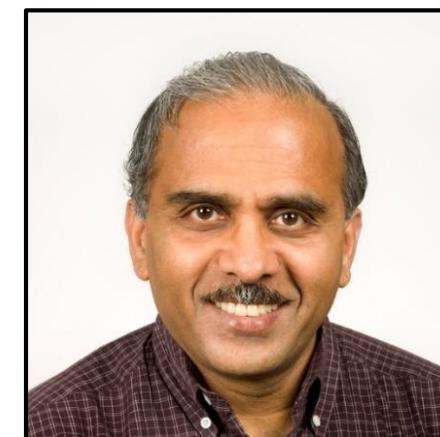
Yang Zou, CMU/NVIDIA



Xiaodong Yang, NVIDIA



Zhidong Yu, NVIDIA

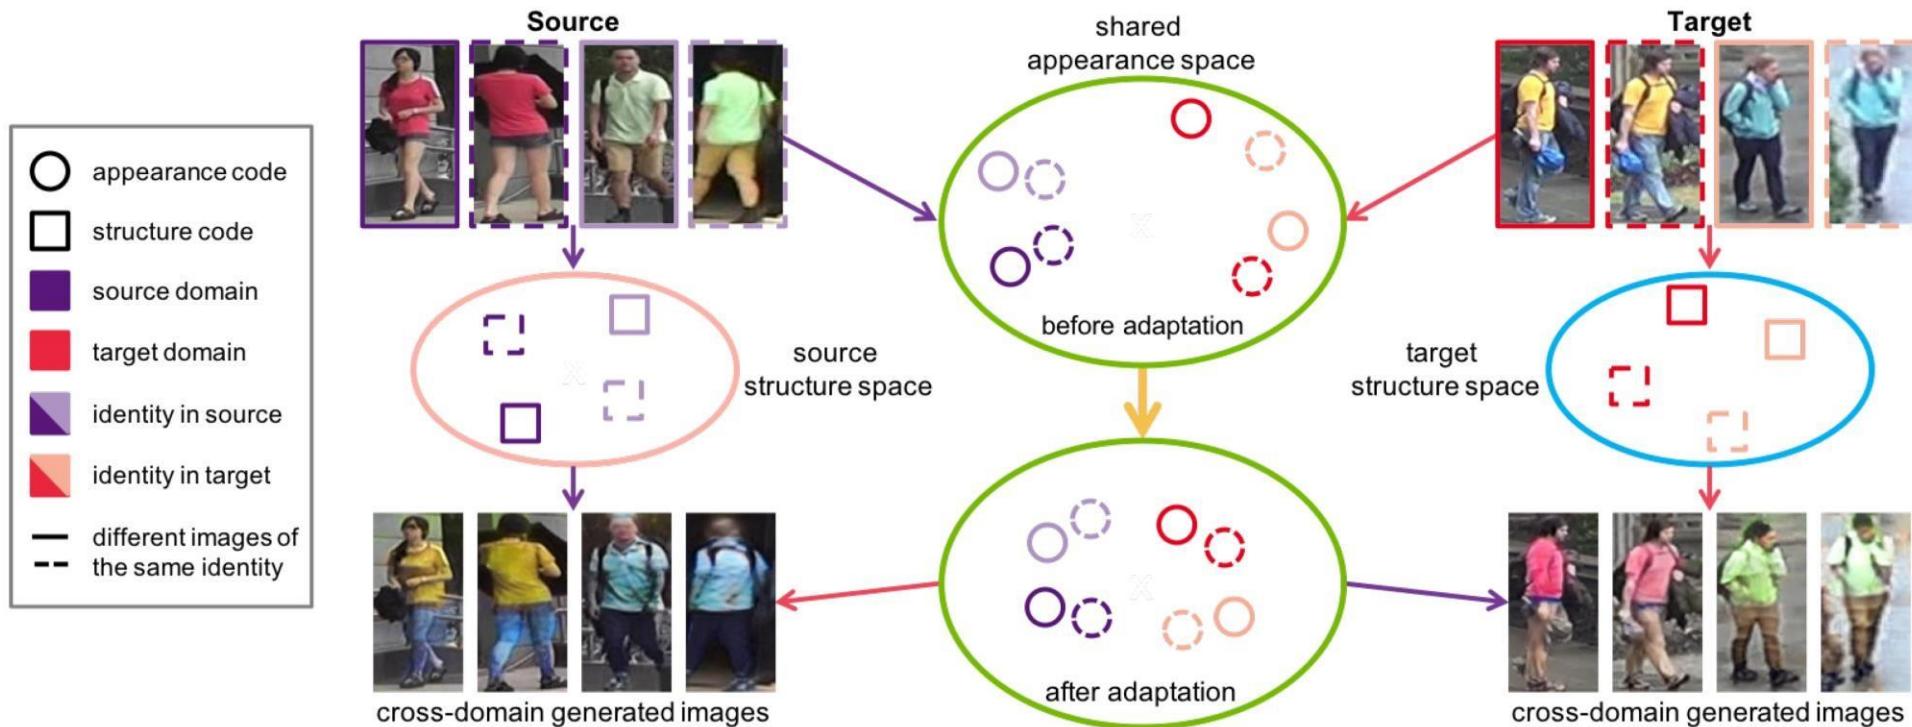


Vijayakumar Bhagavatula,
CMU

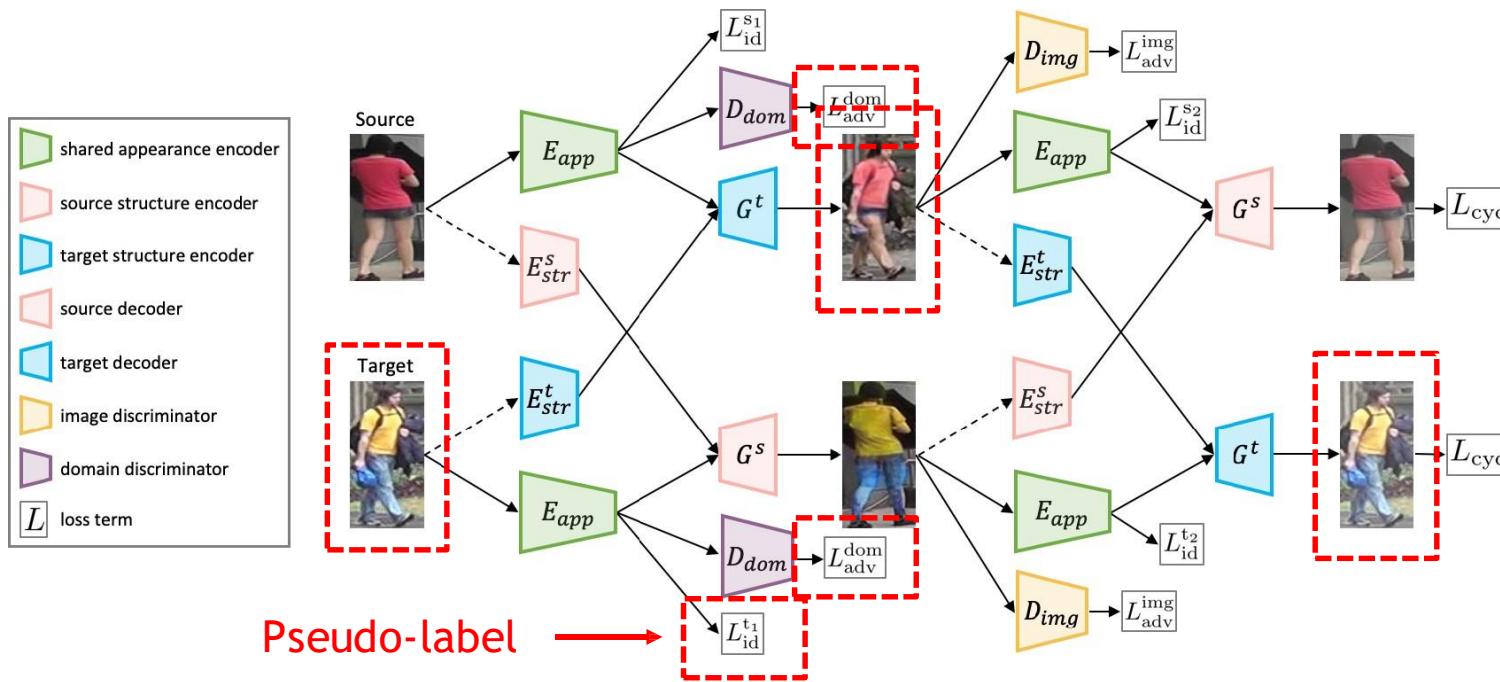


Jan Kautz, NVIDIA

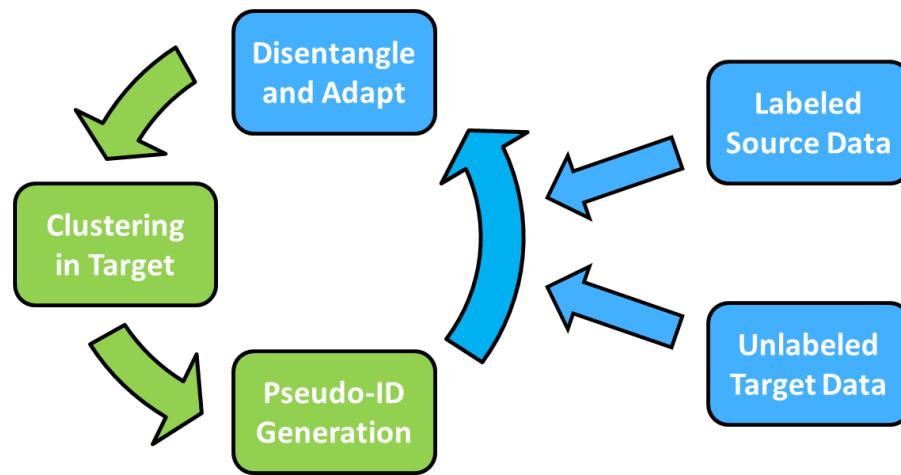
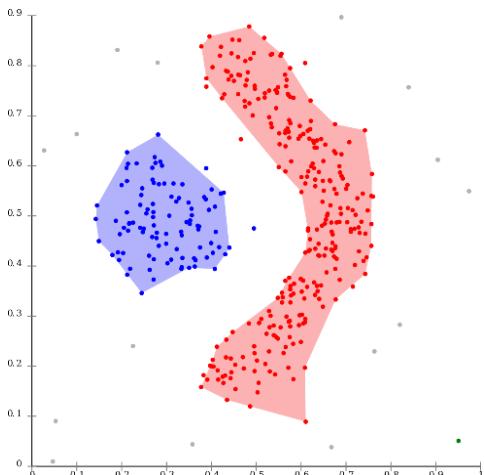
Person Re-ID as An Open Set UDA Problem



DG-Net++: Joint Disentangling and Adaptation



Pseudo-label



- Cycle consistency as self-supervision to disentangle shared representation and reduce domain gap
- Adversarial loss to enforce global distribution matching
- DBSCAN as a clustering-based proxy-label approach to give pseudo-label
- Iterative disentangle & self-training

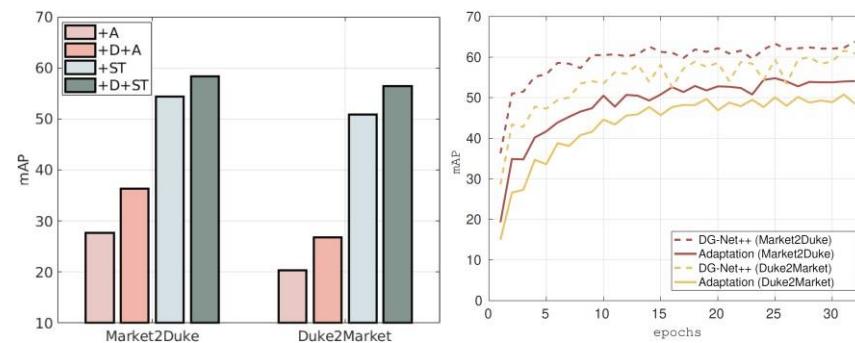
Quantitative Results and Ablation Studies

| Methods | Market-1501 → DukeMTMC-reID | | | | DukeMTMC-reID → Market-1501 | | | |
|------------|-----------------------------|-------------|-------------|-------------|-----------------------------|-------------|-------------|-------------|
| | Rank@1 | Rank@5 | Rank@10 | mAP | Rank@1 | Rank@5 | Rank@10 | mAP |
| SPGAN [6] | 41.1 | 56.6 | 63.0 | 22.3 | 51.5 | 70.1 | 76.8 | 22.8 |
| AIDL [51] | 44.3 | 59.6 | 65.0 | 23.0 | 58.2 | 74.8 | 81.1 | 26.5 |
| MMFA [31] | 45.3 | 59.8 | 66.3 | 24.7 | 56.7 | 75.0 | 81.8 | 27.4 |
| HHL [64] | 46.9 | 61.0 | 66.7 | 27.2 | 62.2 | 78.8 | 84.0 | 31.4 |
| CAL [38] | 55.4 | - | - | 36.7 | 64.3 | - | - | 34.5 |
| ARN [30] | 60.2 | 73.9 | 79.5 | 33.4 | 70.3 | 80.4 | 86.3 | 39.4 |
| ECN [65] | 63.3 | 75.8 | 80.4 | 40.4 | 75.1 | 87.6 | 91.6 | 43.0 |
| PDA [29] | 63.2 | 77.0 | 82.5 | 45.1 | 75.2 | 86.3 | 90.2 | 47.6 |
| CR-GAN [4] | 68.9 | 80.2 | 84.7 | 48.6 | 77.7 | 89.7 | 92.7 | 54.0 |
| IPL [42] | 68.4 | 80.1 | 83.5 | 49.0 | 75.8 | 89.5 | 93.2 | 53.7 |
| SSG [11] | 73.0 | 80.6 | 83.2 | 53.4 | 80.0 | 90.0 | 92.4 | 58.3 |
| DG-Net++ | 78.9 | 87.8 | 90.4 | 63.8 | 82.1 | 90.2 | 92.7 | 61.7 |

| Methods | Market-1501 → MSMT17 | | | DukeMTMC-reID → MSMT17 | | | mAP | |
|------------|----------------------|-------------|-------------|------------------------|-------------|-------------|-------------|-------------|
| | Rank@1 | Rank@5 | Rank@10 | mAP | Rank@1 | Rank@5 | | |
| PTGAN [52] | 10.2 | - | 24.4 | 2.9 | 11.8 | - | 27.4 | 3.3 |
| ENC [65] | 25.3 | 36.3 | 42.1 | 8.5 | 30.2 | 41.5 | 46.8 | 10.2 |
| SSG [11] | 31.6 | - | 49.6 | 13.2 | 32.2 | - | 51.2 | 13.3 |
| DG-Net++ | 48.4 | 60.9 | 66.1 | 22.1 | 48.8 | 60.9 | 65.9 | 22.1 |

| Methods | MSMT17 → Market-1501 | | | MSMT17 → DukeMTMC-reID | | | mAP | |
|-----------|----------------------|-------------|-------------|------------------------|-------------|-------------|-------------|-------------|
| | Rank@1 | Rank@5 | Rank@10 | mAP | Rank@1 | Rank@5 | | |
| PAUL [55] | 68.5 | - | - | 40.1 | 72.0 | - | - | 53.2 |
| DG-Net++ | 83.1 | 91.5 | 94.3 | 64.6 | 75.2 | 73.6 | 86.9 | 58.2 |

| Methods | Market-1501 → DukeMTMC-reID | | | | DukeMTMC-reID → Market-1501 | | | |
|----------|-----------------------------|-------------|-------------|-------------|-----------------------------|-------------|-------------|-------------|
| | Rank@1 | Rank@5 | Rank@10 | mAP | Rank@1 | Rank@5 | Rank@10 | mAP |
| Baseline | 37.4 | 52.4 | 58.4 | 19.3 | 39.7 | 57.9 | 64.3 | 15.0 |
| +A+ST | 71.4 | 81.8 | 85.7 | 57.5 | 75.7 | 86.4 | 90.1 | 57.1 |
| +D | 44.5 | 60.6 | 66.7 | 24.2 | 50.1 | 68.0 | 73.9 | 26.8 |
| +D+A | 53.2 | 68.7 | 73.8 | 36.3 | 52.2 | 70.7 | 77.0 | 28.6 |
| +D+ST | 74.2 | 82.8 | 86.5 | 58.4 | 78.0 | 87.1 | 90.3 | 56.5 |
| +D+A+ST | 78.9 | 87.8 | 90.4 | 63.8 | 82.1 | 90.2 | 92.7 | 61.7 |



| Method | Metric | Market→Duke | Duke→Market | MSMT→Market | Market→MSMT | MSMT→Duke | Duke→MSMT |
|-------------|--------|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|
| DG-Net [59] | Rank@1 | 42.6 | 56.1 | 61.8 | 17.1 | 61.9 | 20.6 |
| | mAP | 24.3 | 26.8 | 33.6 | 5.4 | 40.7 | 6.4 |
| DG-Net++ | Rank@1 | 78.9 (+36.3) | 82.1 (+26.0) | 83.1 (+21.3) | 48.4 (+31.3) | 75.2 (+13.3) | 48.8 (+28.2) |
| | mAP | 63.8 (+39.5) | 61.7 (+34.9) | 64.6 (+31.0) | 22.1 (+16.7) | 58.2 (+17.5) | 22.1 (+15.7) |

Qualitative Results

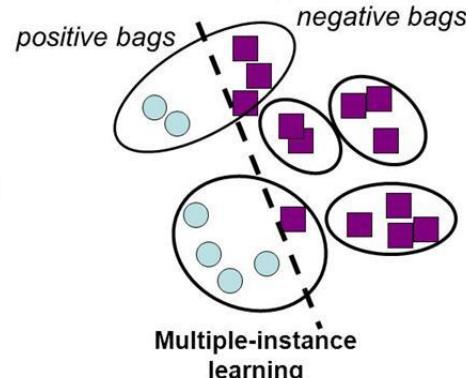
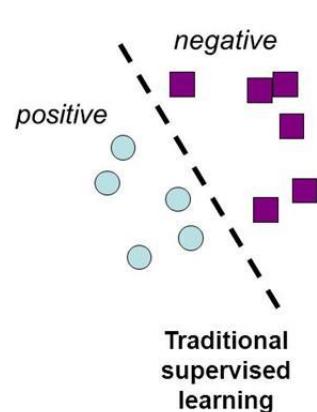
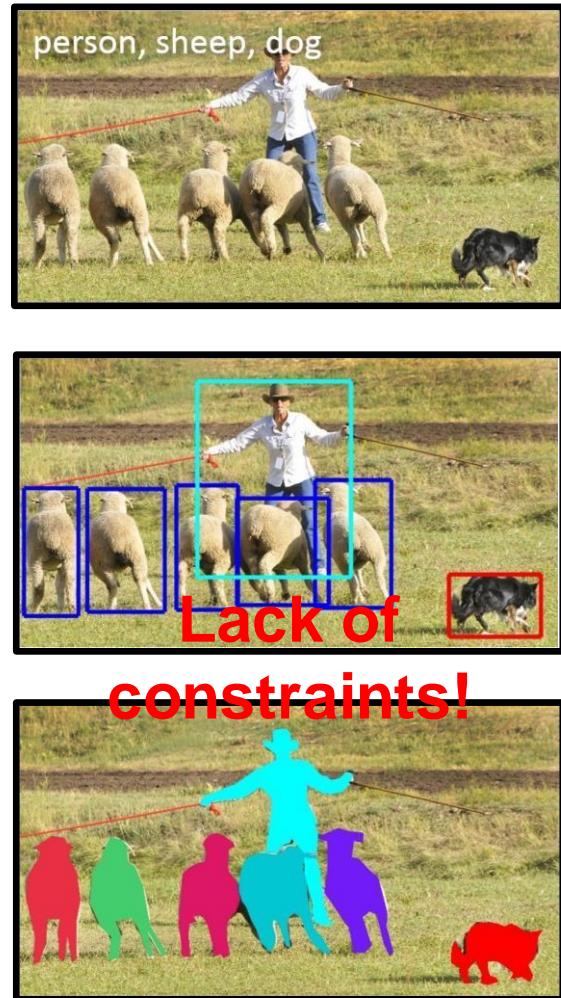




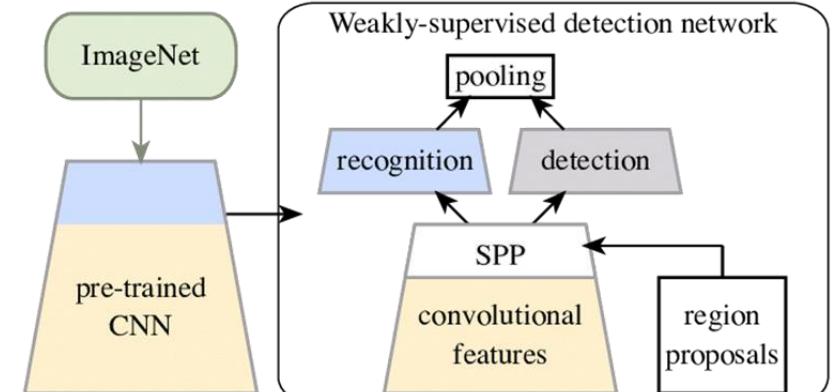
Part II: Unifying Levels of Supervision with Proxy-Label Approaches



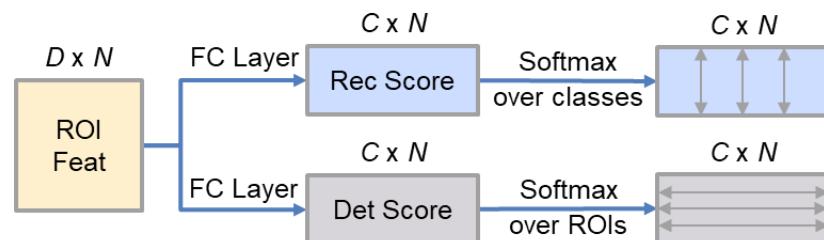
Weakly Supervised Object Detection



Concept of multiple-instance-learning



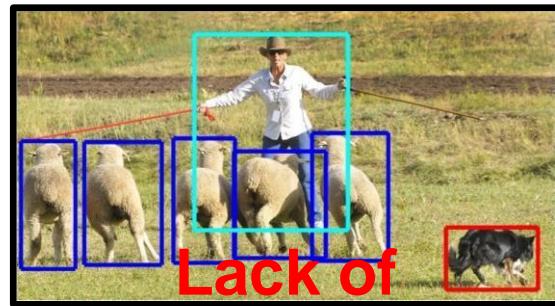
Overview of WSDDN



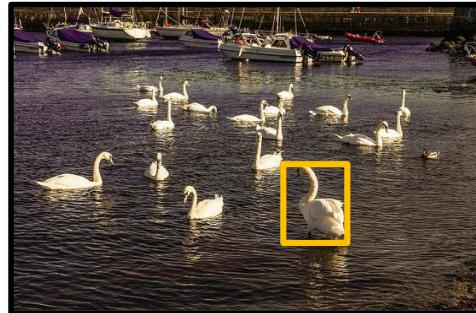
$$\mathcal{L}_{\text{img}}(w) = - \sum_{c \in C} y(c) \log(\phi_w(c))$$

$$\text{ROI Score} \xrightarrow[\text{Sum pooling over ROIs}]{\mathcal{L}_{\text{img}}} \begin{matrix} C \\ \text{Image Score} \end{matrix} \xrightarrow{\text{Image Label}}$$

Weakly Supervised Object Detection



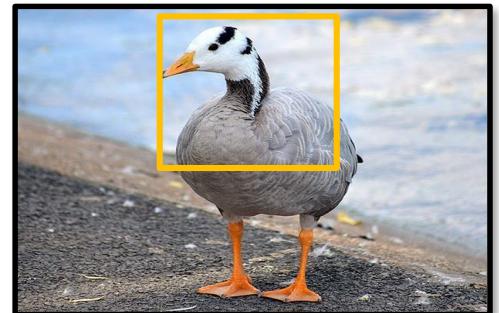
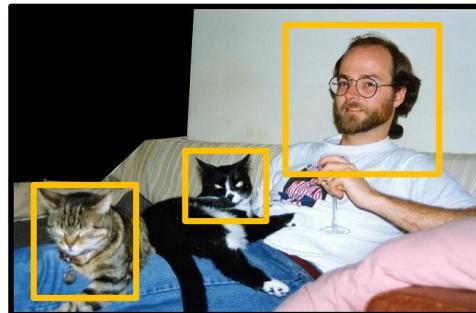
Missing Instance



Grouped Instance

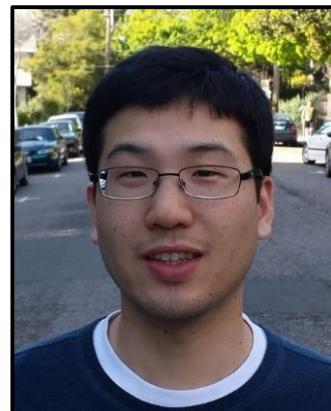
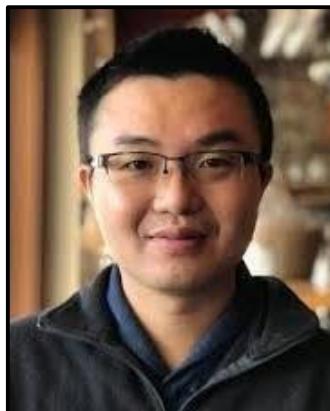
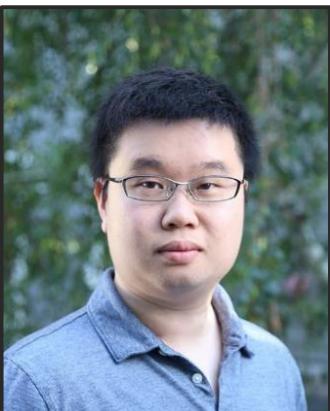


Part Domination



Instance-aware, Context-focused, and Memory-efficient Weakly Supervised Object Detection (CVPR20)

UFO²: A Unified Framework towards Omni-supervised Object Detection (ECCV20)



Zhongzheng Ren
UIUC/NVIDIA

Zhiding Yu
NVIDIA

Xiaodong Yang
NVIDIA

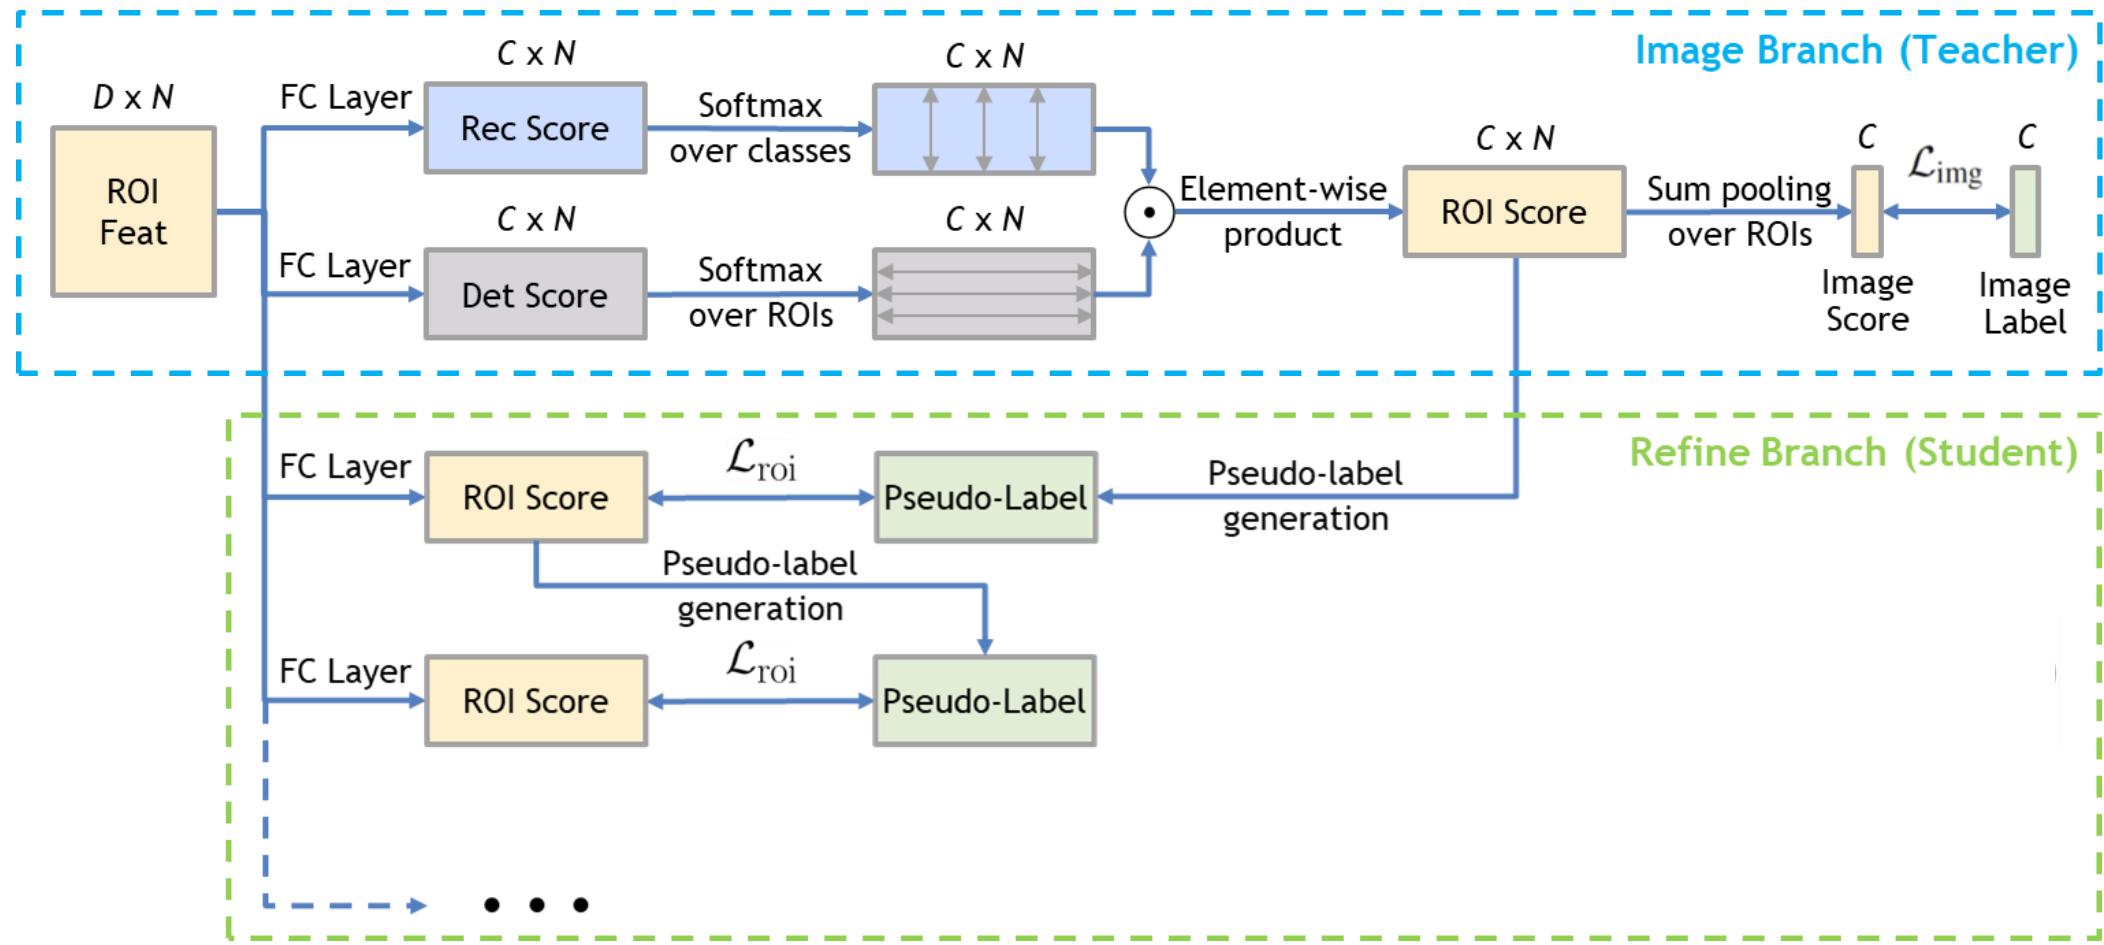
Ming-Yu Liu
NVIDIA

Yong Jae Lee
UC Davis

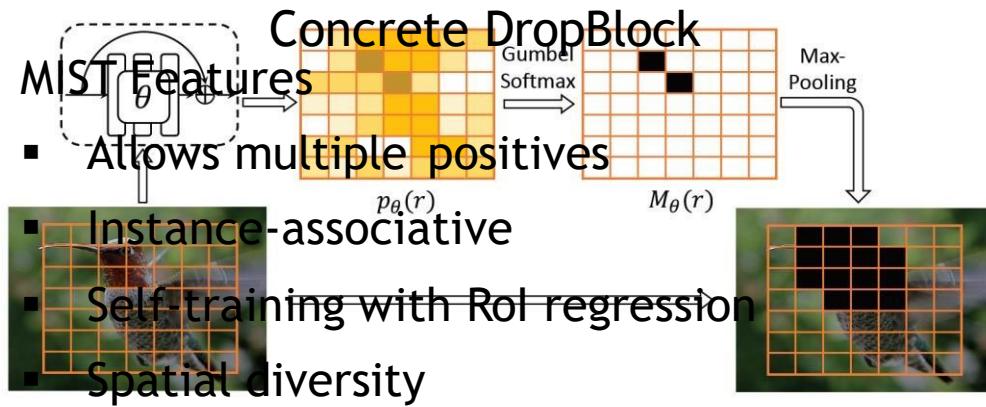
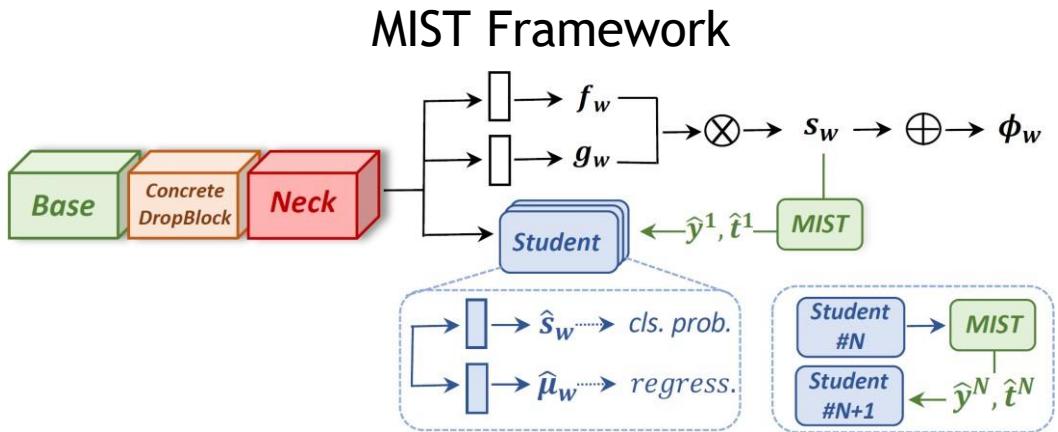
Alex Schwing
UIUC

Jan Kautz
NVIDIA

Online Instance Classifier Refinement (OICR)



Multiple Instance Self-Training (MIST)

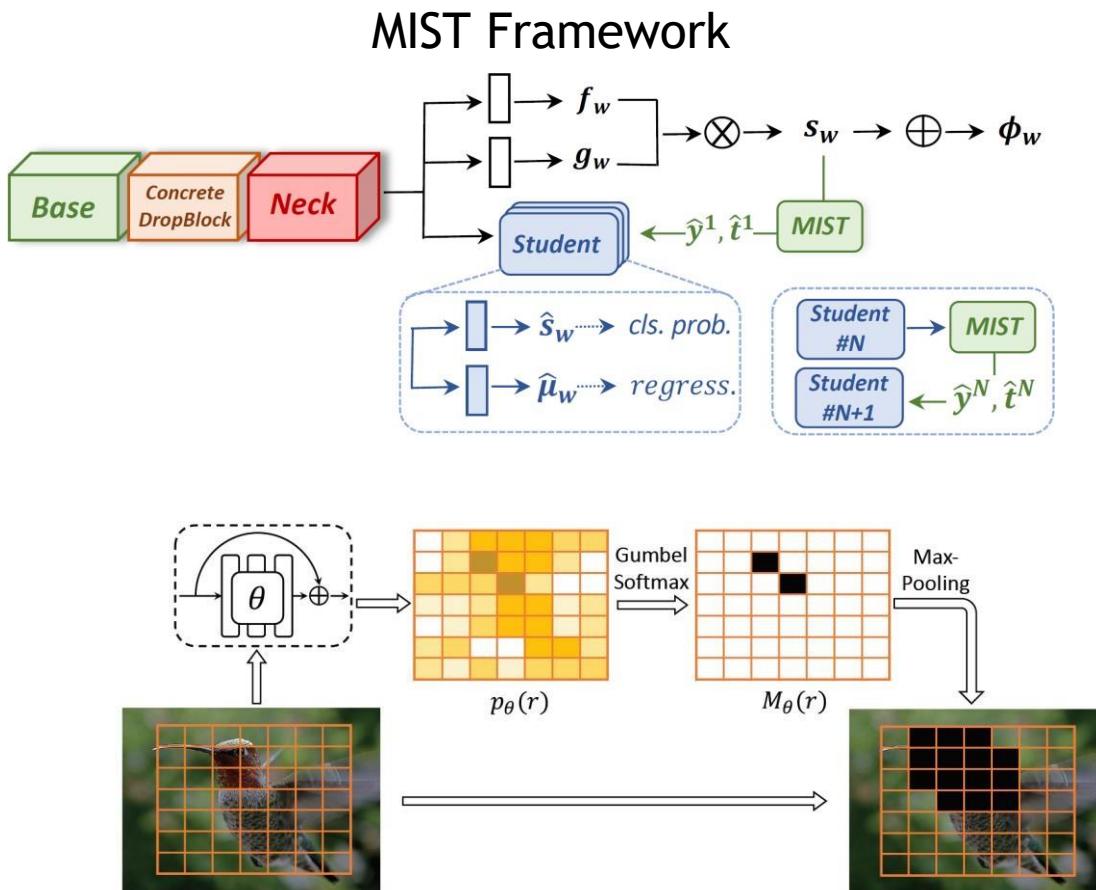


| Weakly Super Metrics | 07 trainval CorLoc | 07 test Det. | 12 trainval CorLoc | 12 test Det. |
|-----------------------|--------------------|--------------|--------------------|--------------|
| Baseline [45]* | 60.8 | 42.5 | - | - |
| Leader + PCL [44] | 62.7 | 43.5 | 63.2 | 40.6 |
| TREND + MIST w/o Reg. | 62.9 | 48.3 | 65.1 | - |
| + MIST | 64.9 | 51.4 | 66.7 | - |
| + Img Spa.-Dropout | 64.3 | 51.1 | 65.9 | - |
| + ROI Spa.-Dropout | 66.8 | 52.4 | 67.3 | - |
| + DropBlock [14] | 67.1 | 52.9 | 68.4 | - |
| + Concrete DropBlock | 68.8 | 54.9 | 70.9 | 52.1 |

Charades Spatial Prior Activity Driven Weakly Supervised Object Detection

| Methods | Val-AP | Val-AP ₅₀ | Test-AP | Test-AP ₅₀ |
|-----------------------------|-------------|----------------------|-------------|-----------------------|
| Fast R-CNN | 18.9 | 38.6 | 19.3 | 39.3 |
| Faster R-CNN | 21.2 | 41.5 | 21.5 | 42.1 |
| WSDDN [5] | - | - | - | 11.5 |
| WCCN [9] | - | - | - | 12.3 |
| PCL [45] | 8.5 | 19.4 | - | - |
| C-MIDN [12] | 9.6 | 21.4 | - | - |
| WSOD2 [61] | 10.8 | 22.7 | - | - |
| Diba <i>et al.</i> [10]+SSD | - | - | - | 13.6 |
| OICR [46]+Ens+FRCNN | 7.7 | 17.4 | - | - |
| Ge <i>et al.</i> [13]+FRCNN | 8.9 | 19.3 | - | - |
| PCL [45]+Ens.+FRCNN | 9.2 | 19.6 | - | - |
| Ours (single-model) | 11.4 | 24.3 | 12.1 | 24.8 |

Multiple Instance Self-Training (MIST)



Algorithm 1 Multiple Instance Self-Training

Input: Image I , class label y , proposals R , threshold τ , percentage p

Output: Pseudo boxes \hat{R}^1

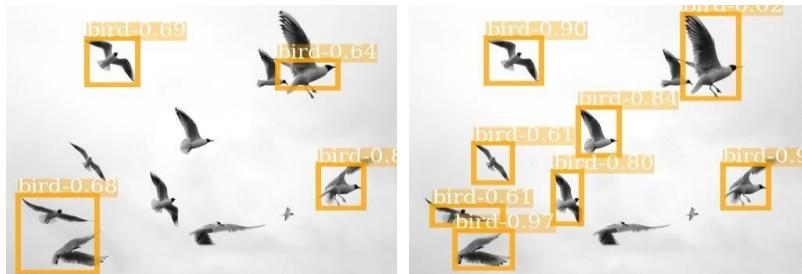
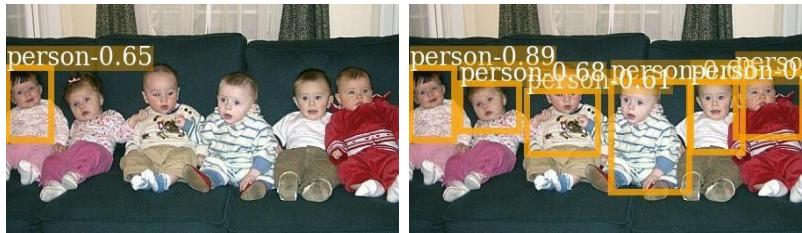
- 1: Feed I into model; get ROI scores s
 - 2: **for** ground-truth class c **do**
 - 3: $R(c)_{sorted} \leftarrow \text{SORT}(s(c, *))$ //sort ROIs by scores of class c
 - 4: $R'(c) \leftarrow \text{top } p\% \text{ percent of } R(c)_{sorted}$
 - 5: $\hat{R}(c) \leftarrow r'_0$ // save first region (top-scoring) $r'_0 \in R'$
 - 6: **for** i in $\{2 \dots |R'(c)|\}$ **do** // start from the second highest
 - 7: APPEND($\hat{R}(c), r'_i$) if $\text{IoU}(r'_i, \hat{r}_j) < \tau, \forall \hat{r}_j \in \hat{R}(c)$
 - 8: **return** $\hat{R}(c)$
-

防止伪标签的检测框在最重要的部位生成，而不是在整个物体上生成，这里采用了concrete DropBlock思路

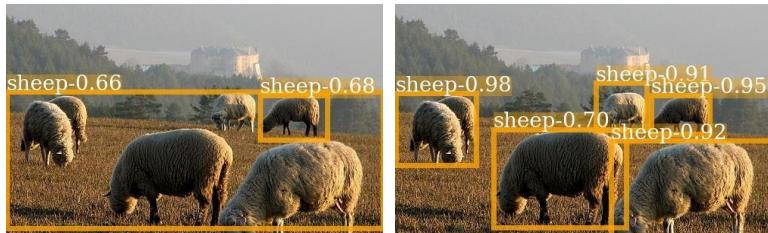
Discriminative parts such as head are zeroed out。

Qualitative Results

Missing Instance



Grouped Instance

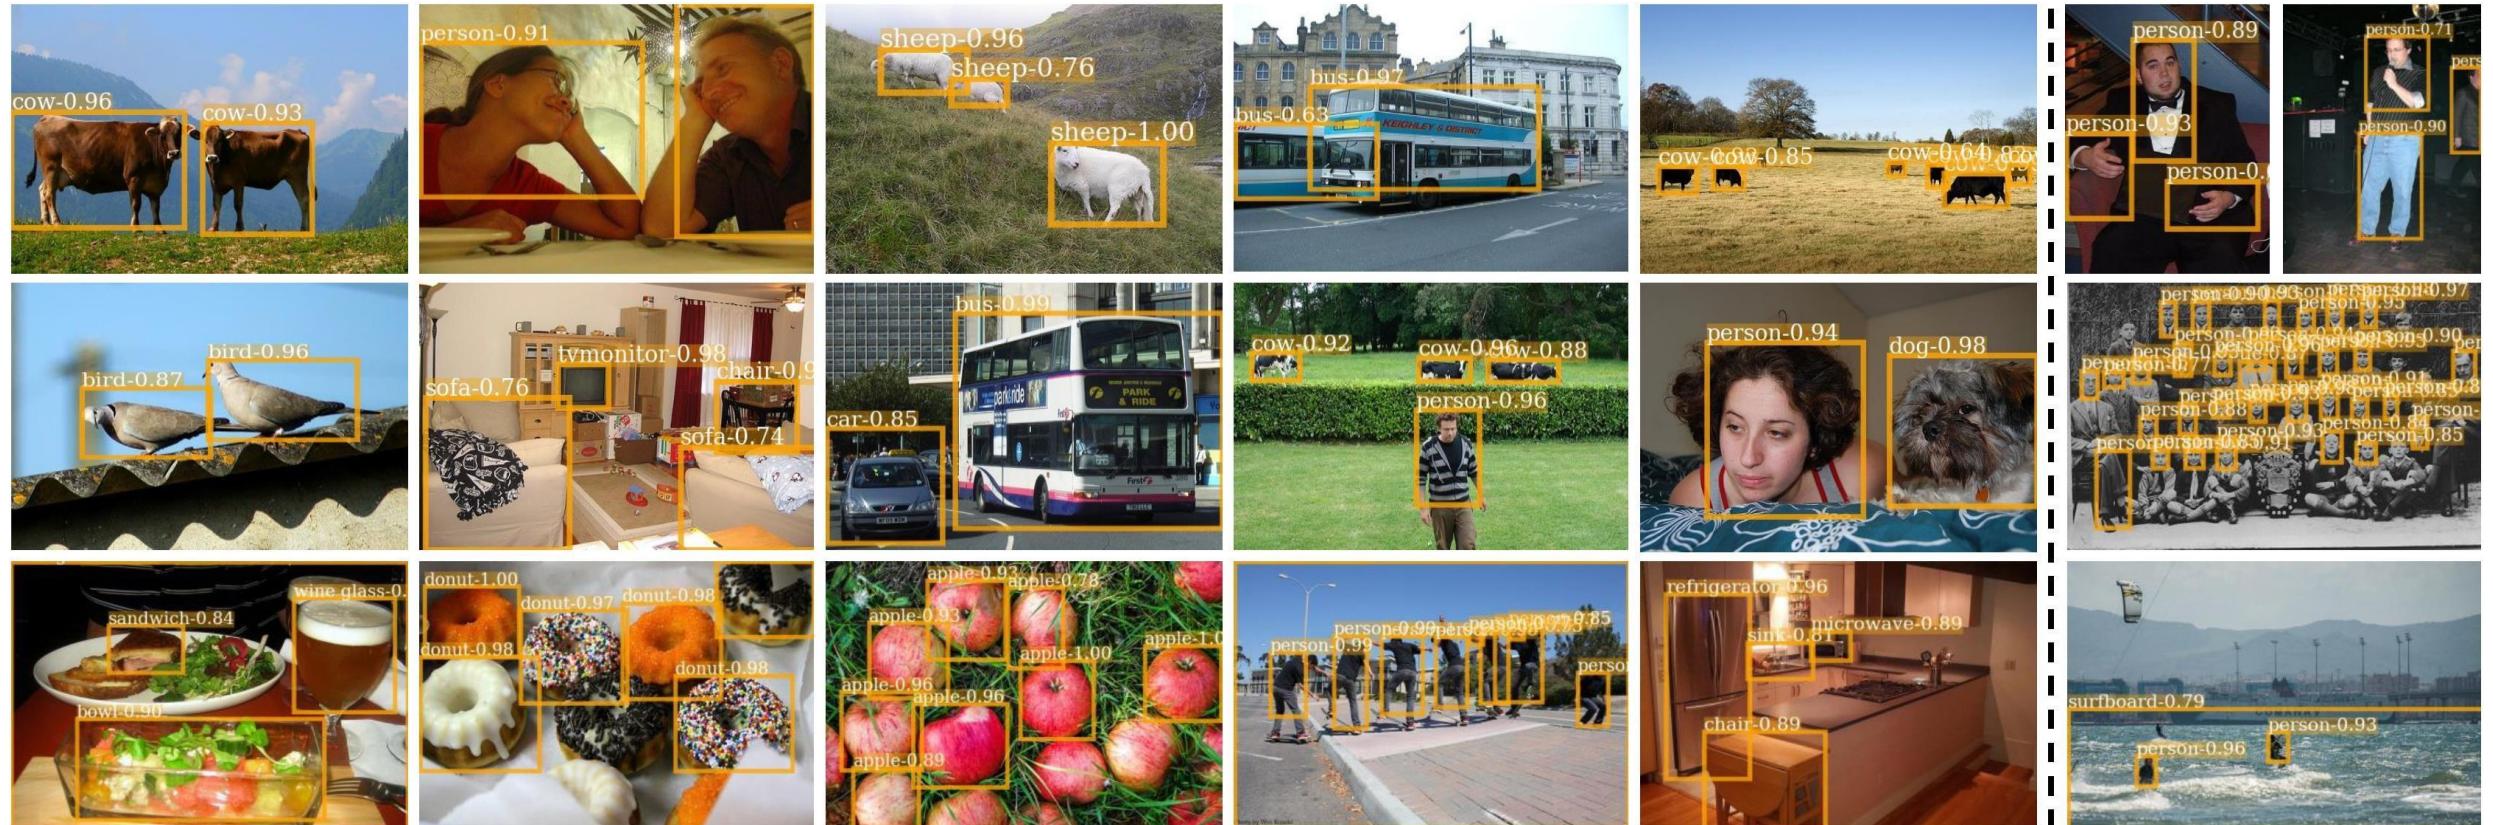


Part Domination



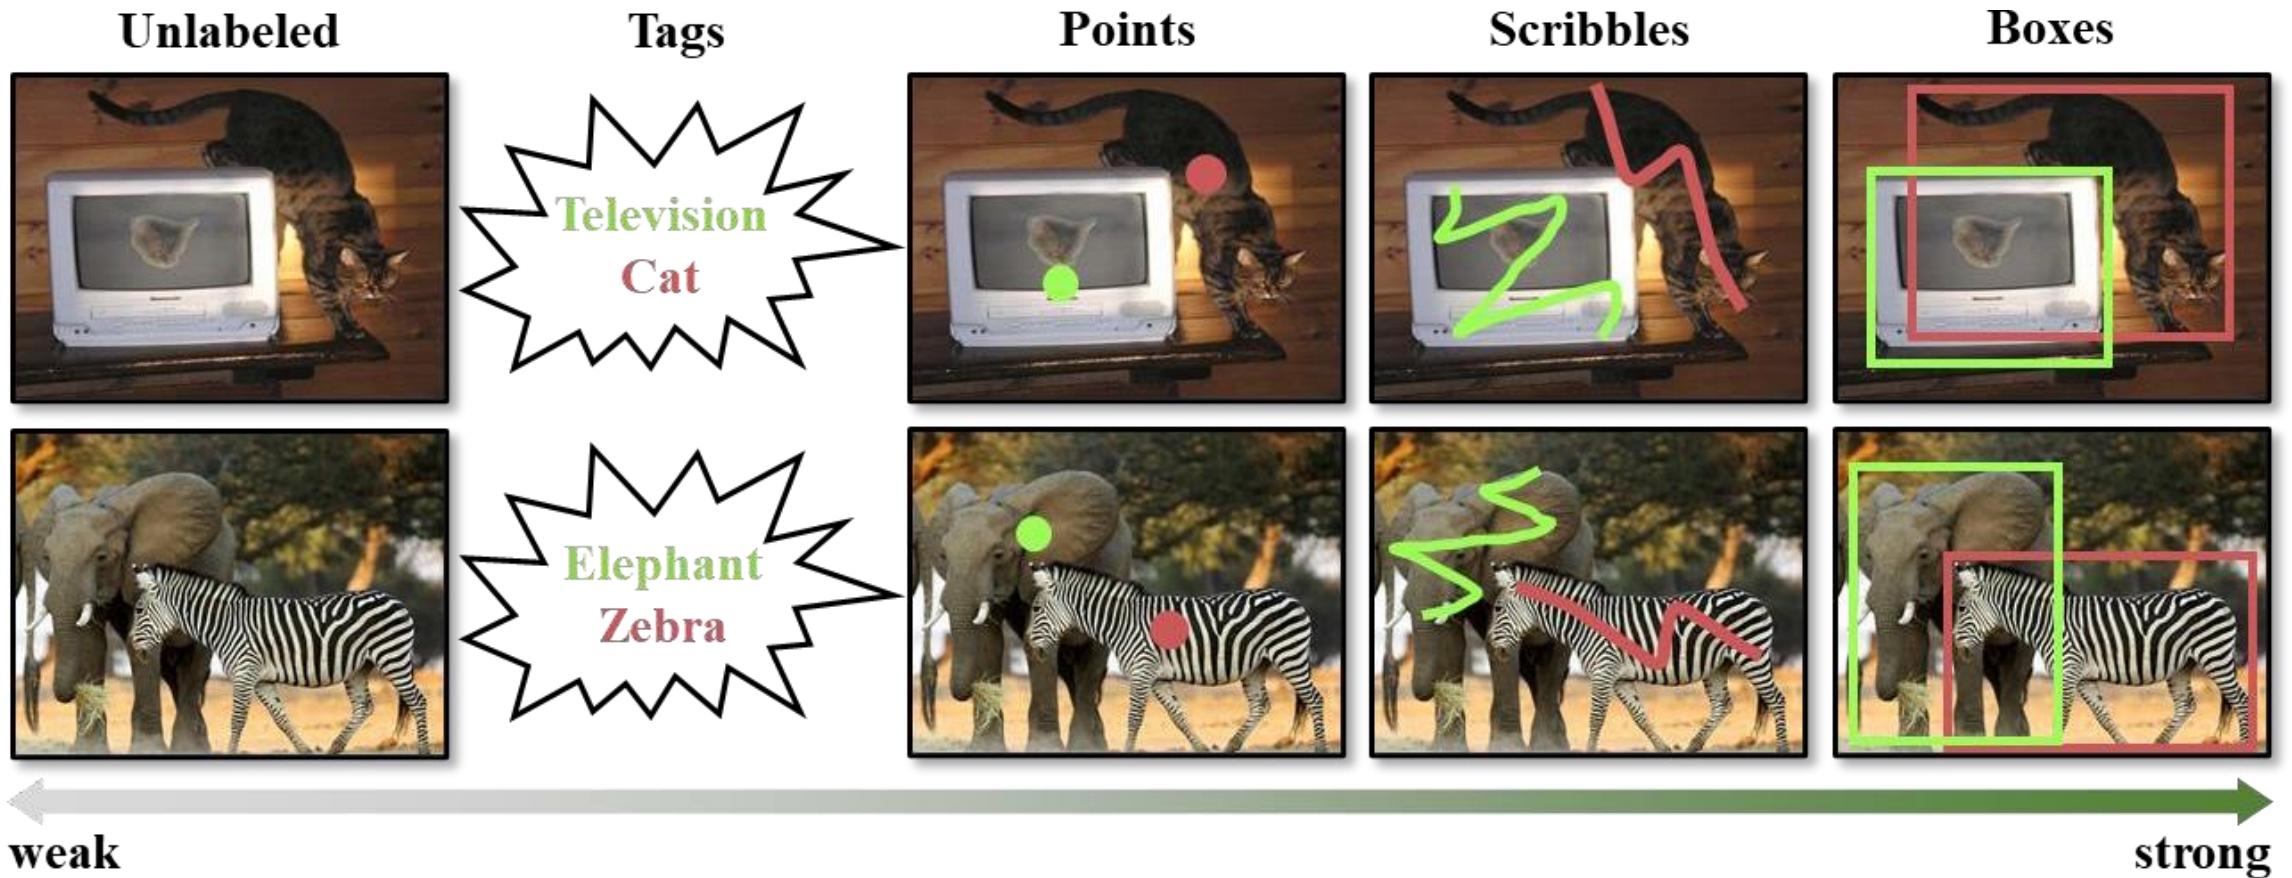
Comparison between our final WSOD model (right in pairs) and OICR (left in pairs)

Qualitative Results

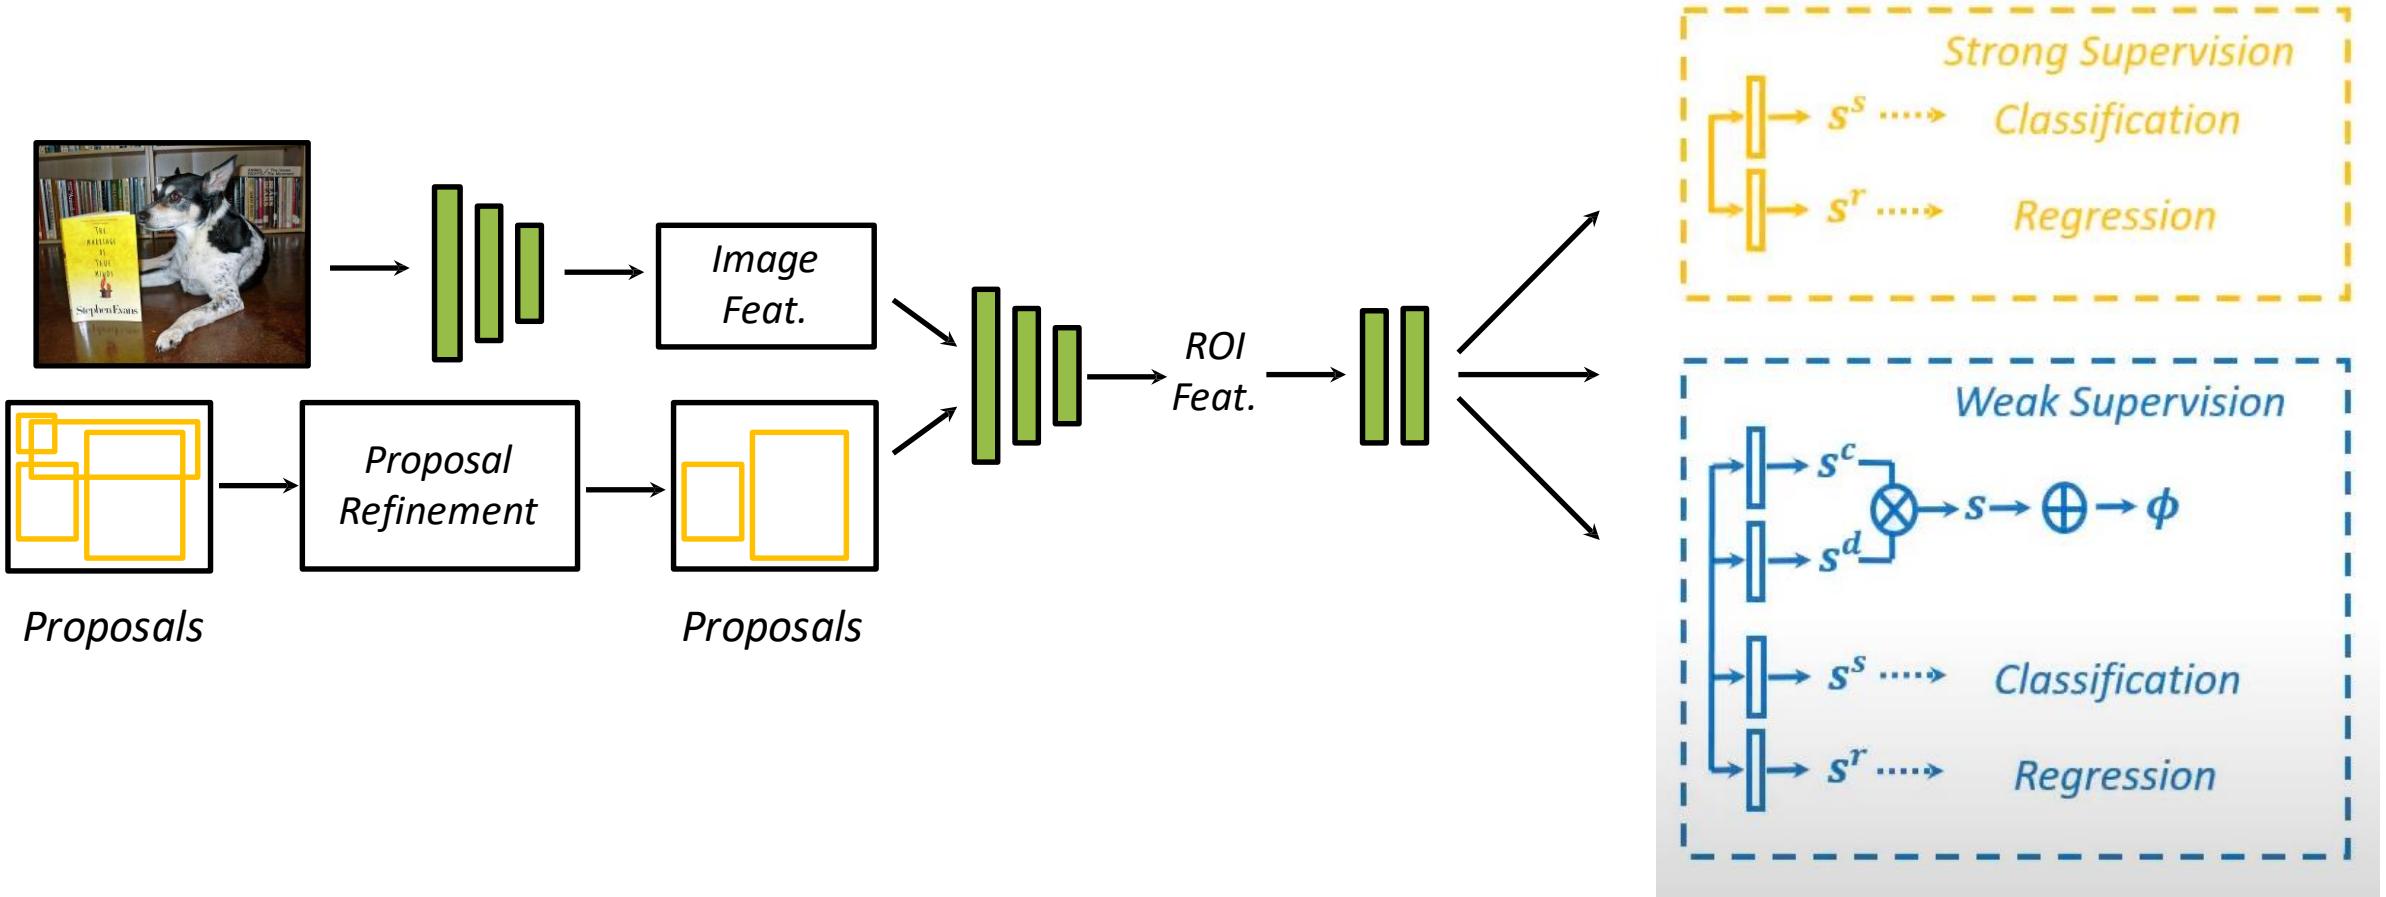


Additional visualizations on different datasets (top: VOC 2007, middle: VOC 2012, bottom: COCO) and examples of failure cases

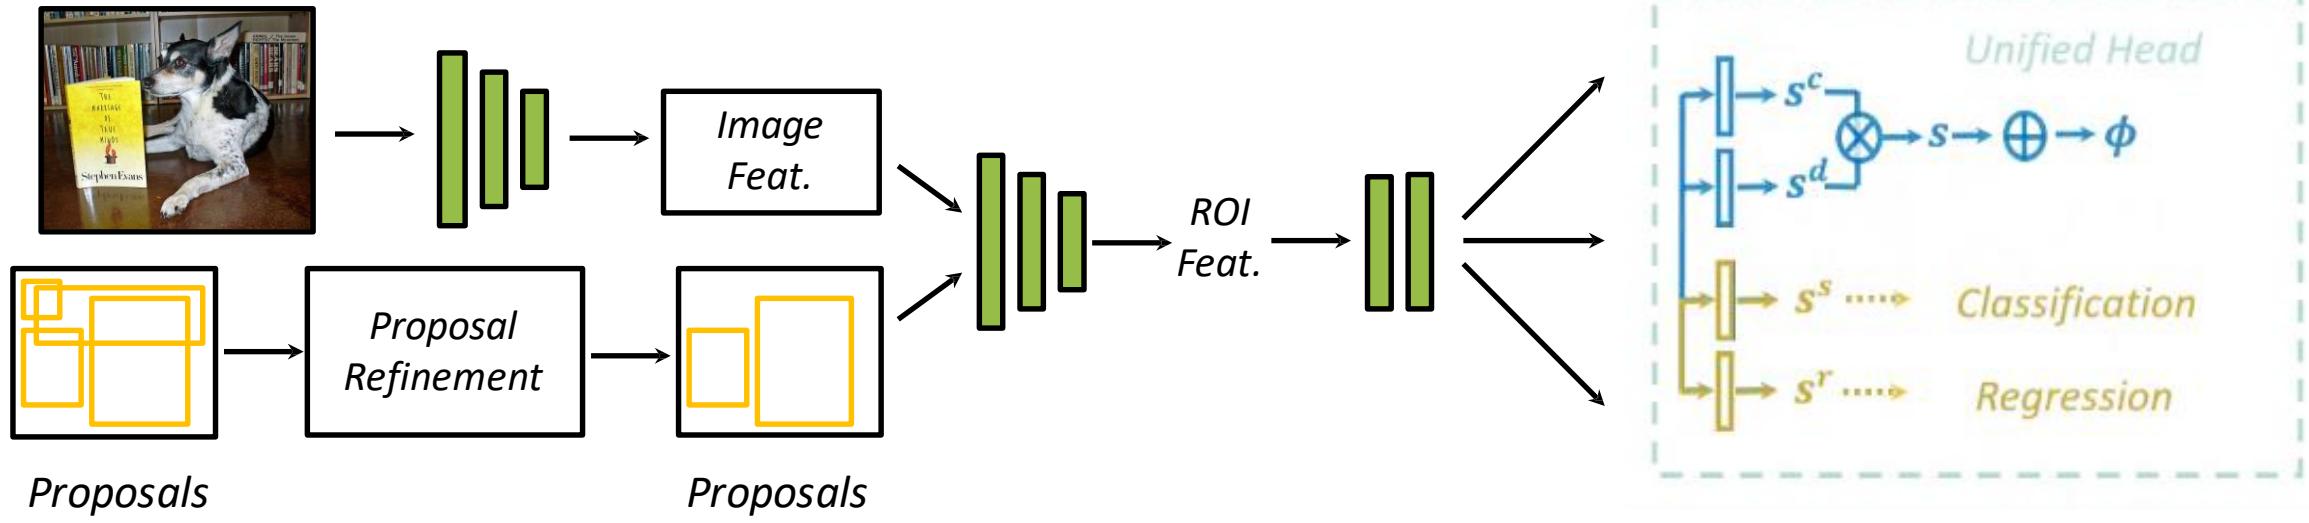
Omni-Supervised Object Detection



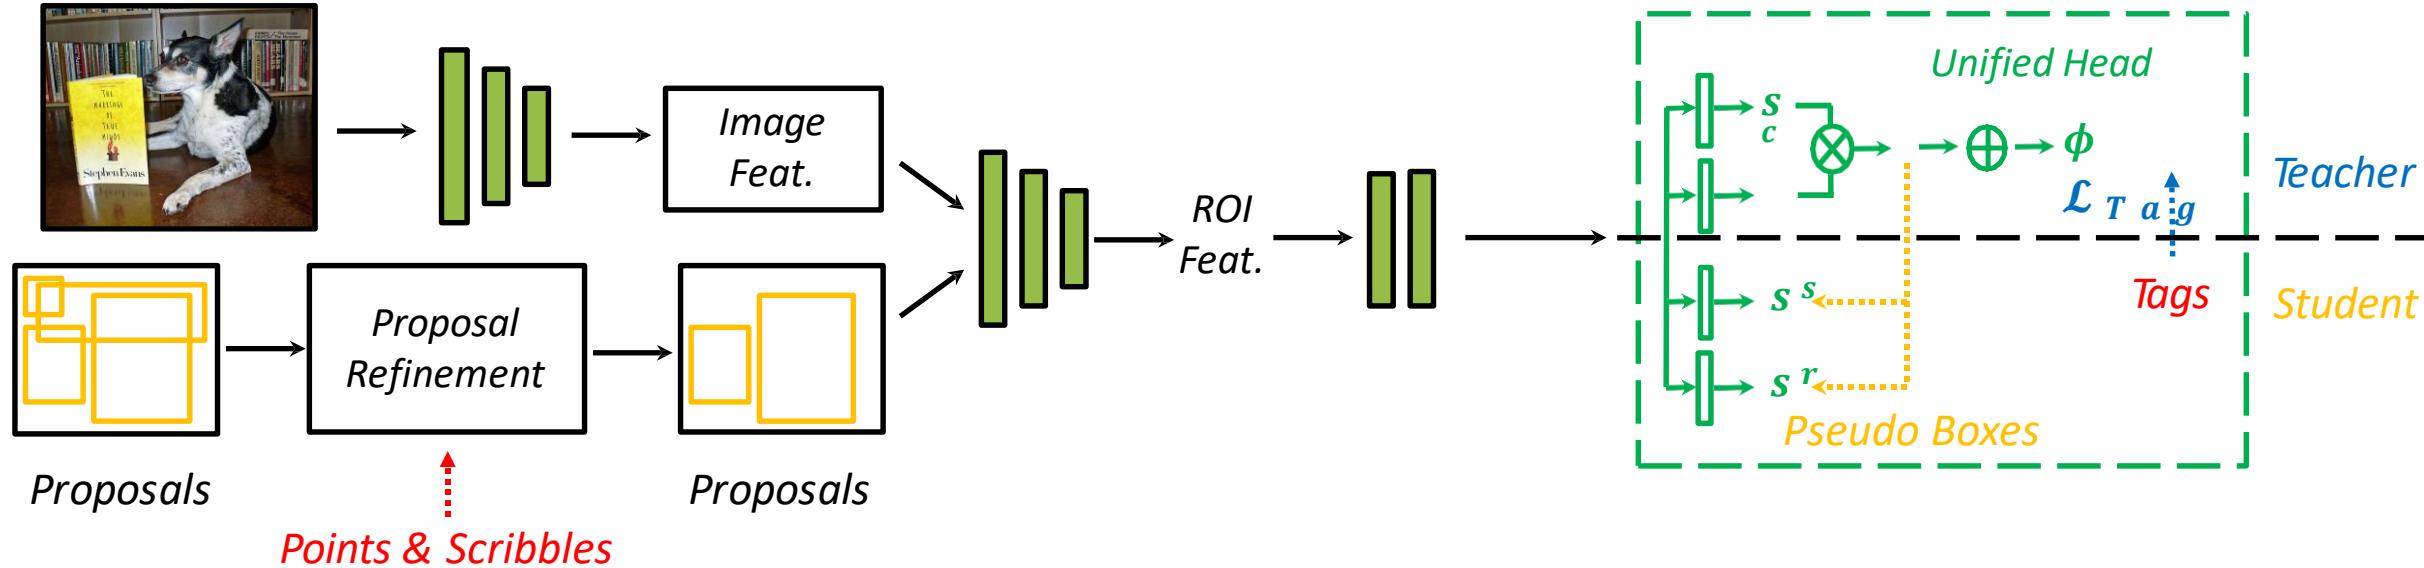
A Unified Omni-Supervised Framework



A Unified Omni-Supervised Framework



Partial Supervision (Tags, Points, Scribbles)



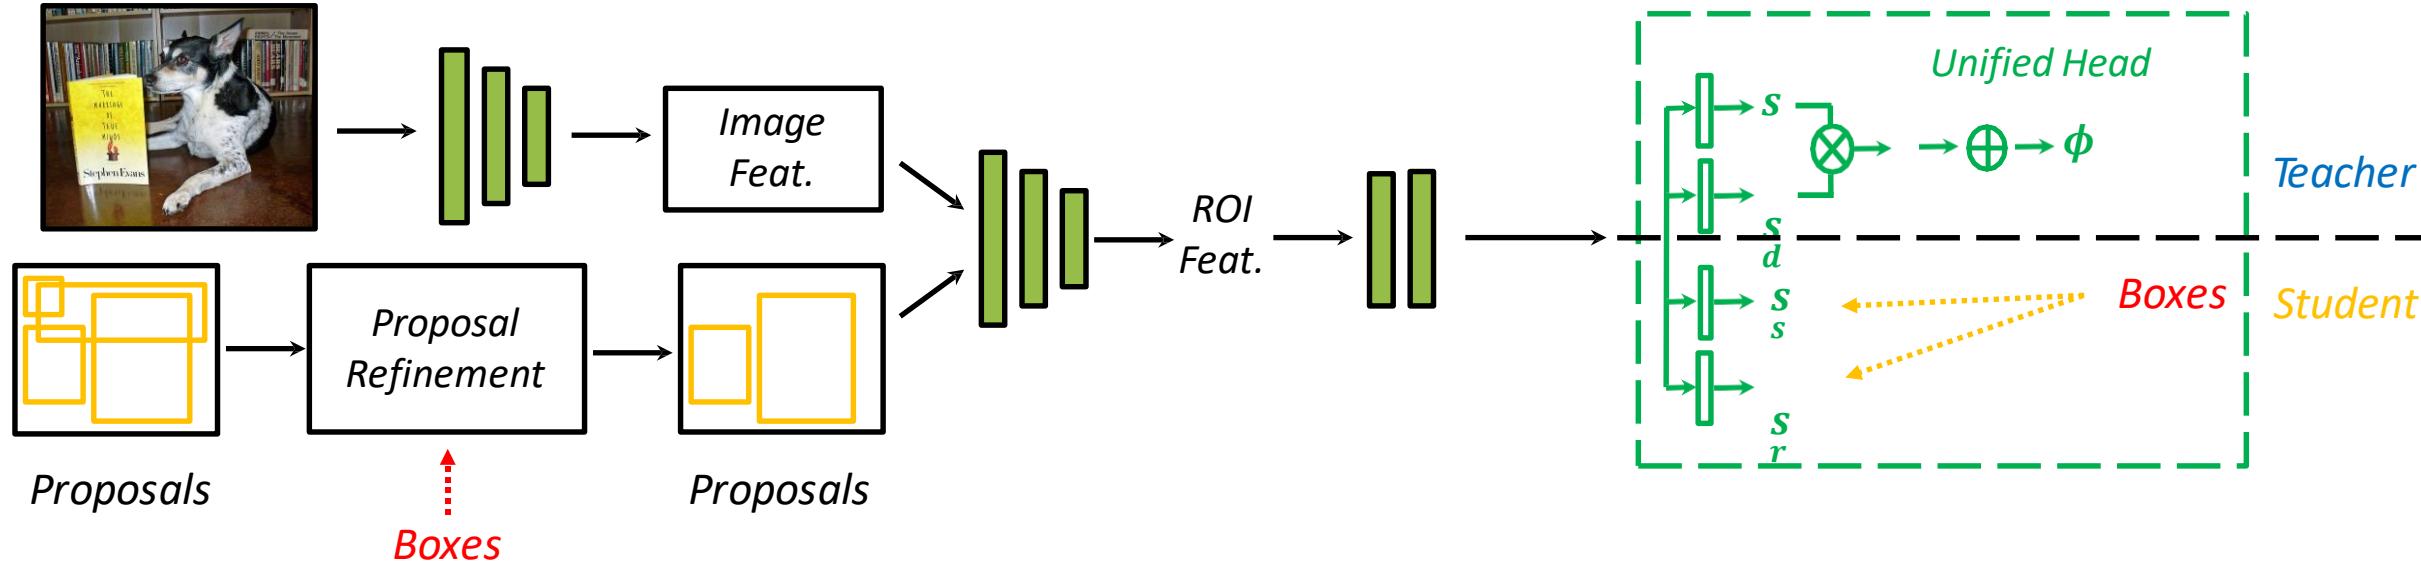
Teacher heads:

- No GT boxes available
- Image-level multi-label classification
(\mathcal{L}_{Tag})

Student heads:

- Generate pseudo-boxes online
- ROI classification
- ROI regression

Strong Supervision (Boxes)

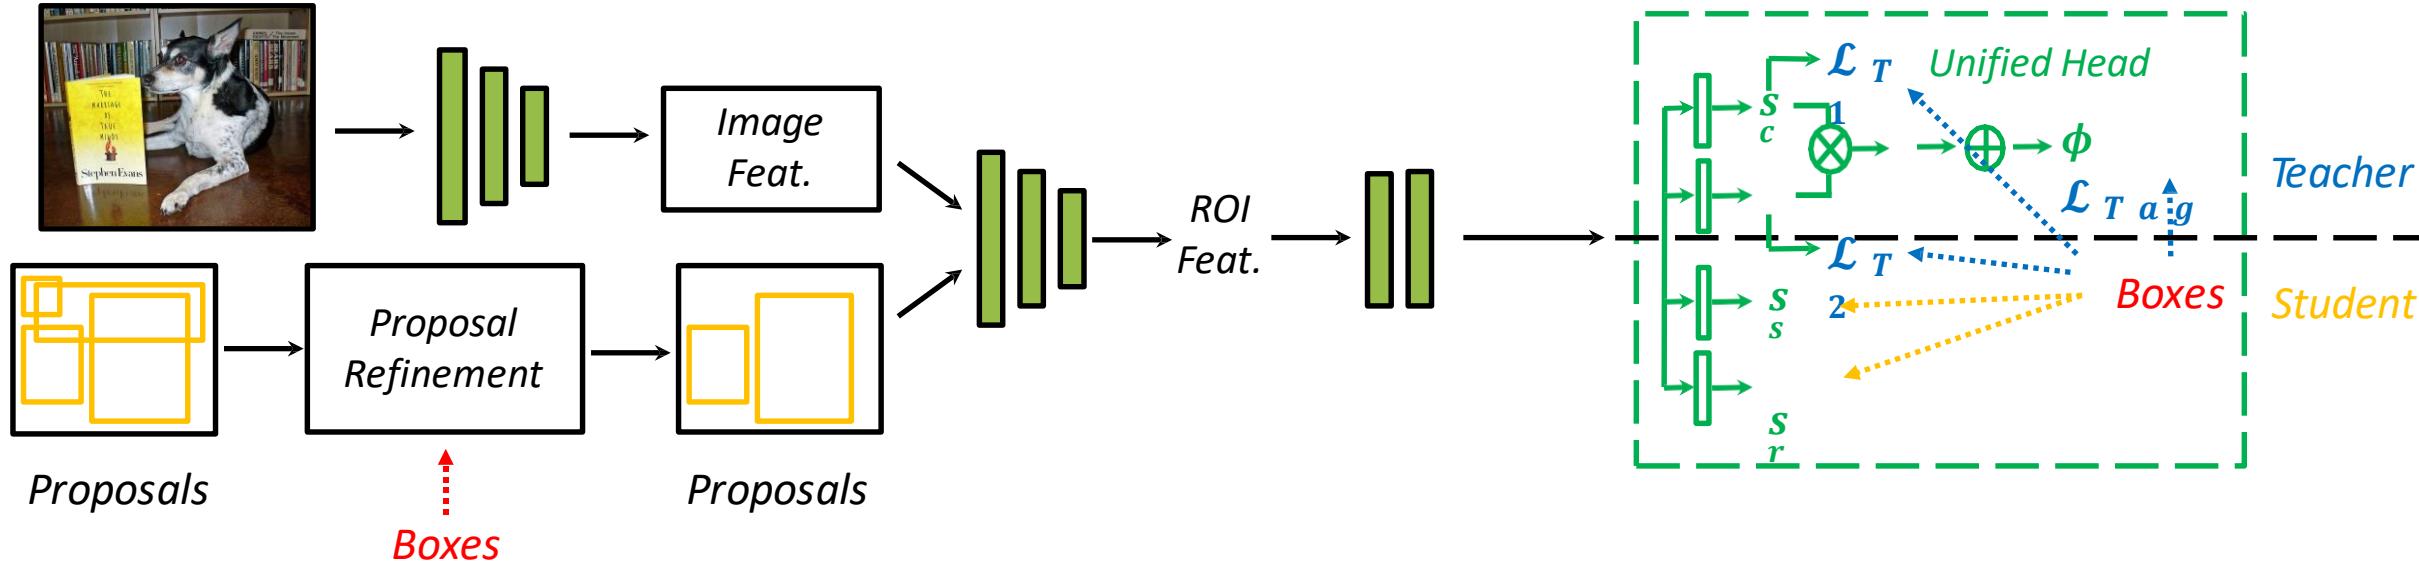


Naïve solution: directly supervise student heads using GT boxes



Issue: Weak Teacher & Strong Students

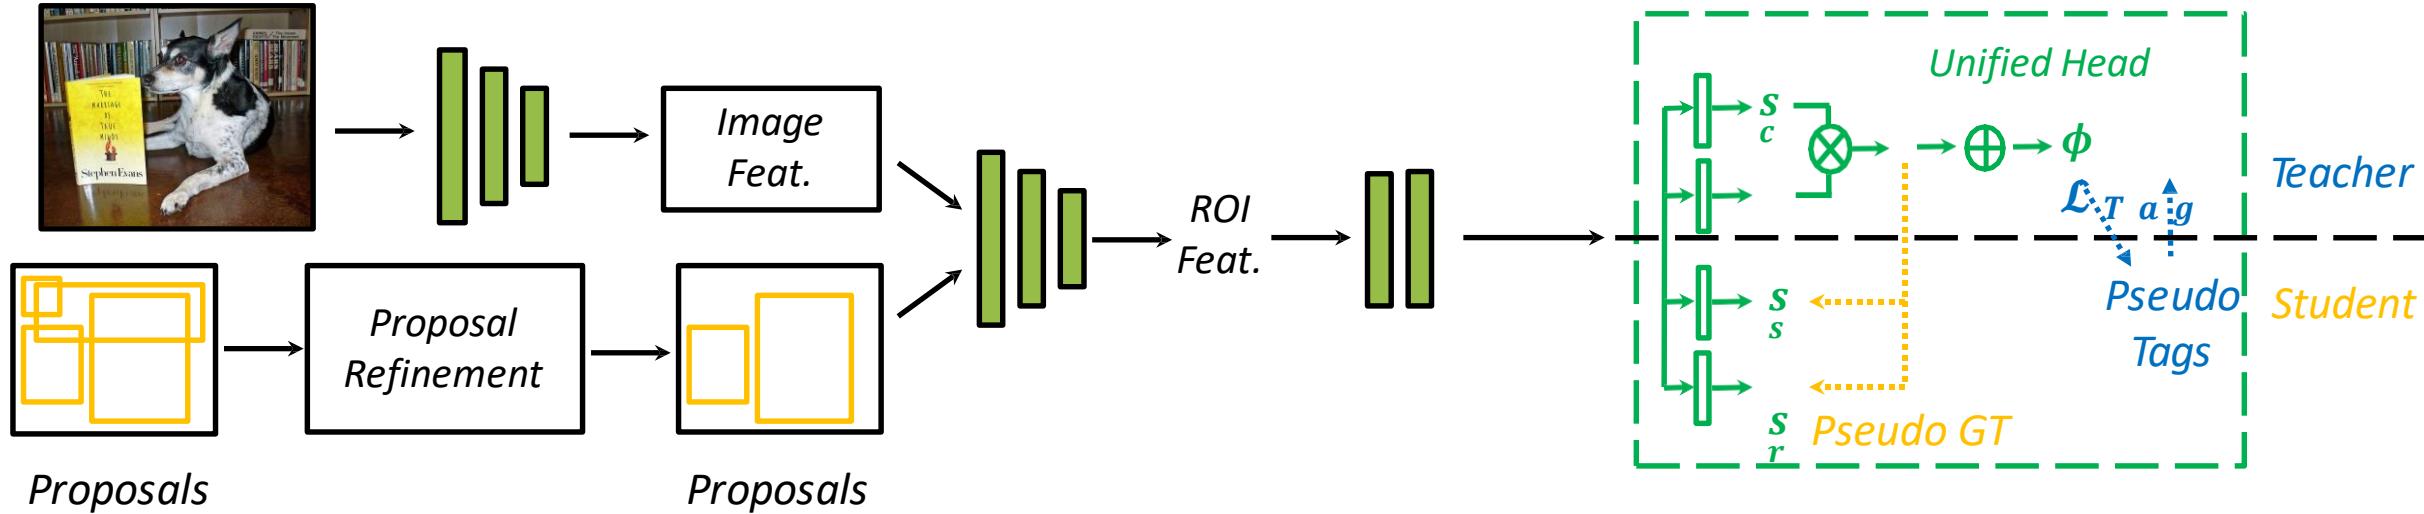
Strong Supervision (Boxes)



Make Teacher **stronger**:

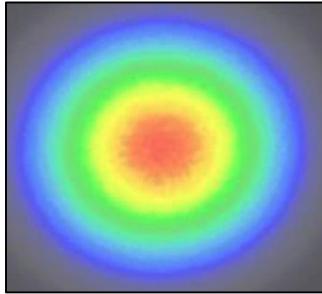
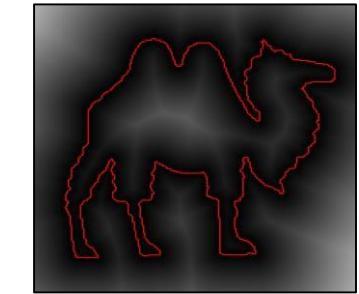
- Image-level multi-label classification ($\mathcal{L}_{T \text{ a } g_s}$)
- ROI classification ($\mathcal{L}_{T \text{ 1}}$)
- ROI objectness regularization ($\mathcal{L}_{T \text{ 2}}$)

Unlabeled Data



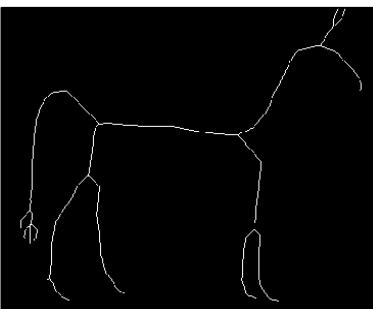
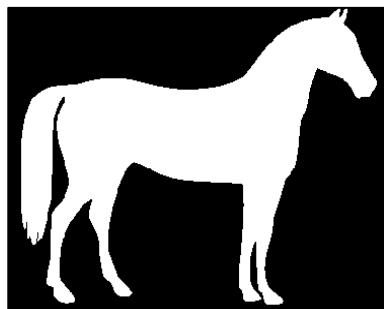
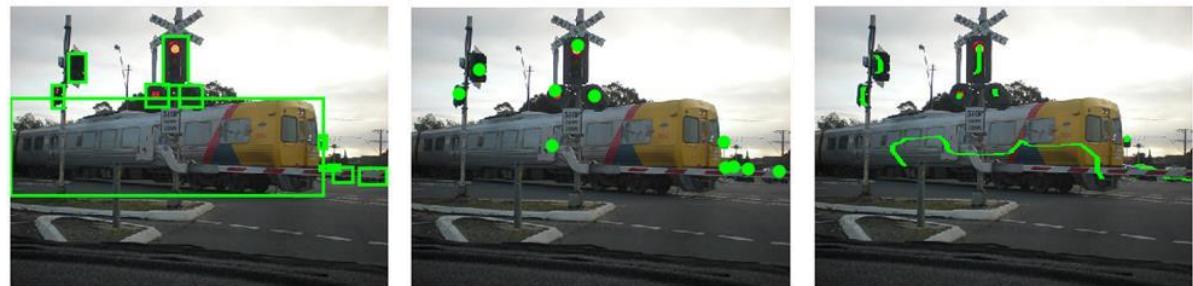
- Generate pseudo-tags by taking confident classes from teacher's predication
- Supervise teacher with pseudo-tags
- Generate pseudo-boxes to supervise student

Partially Labeled MS-COCO



Distance Transform

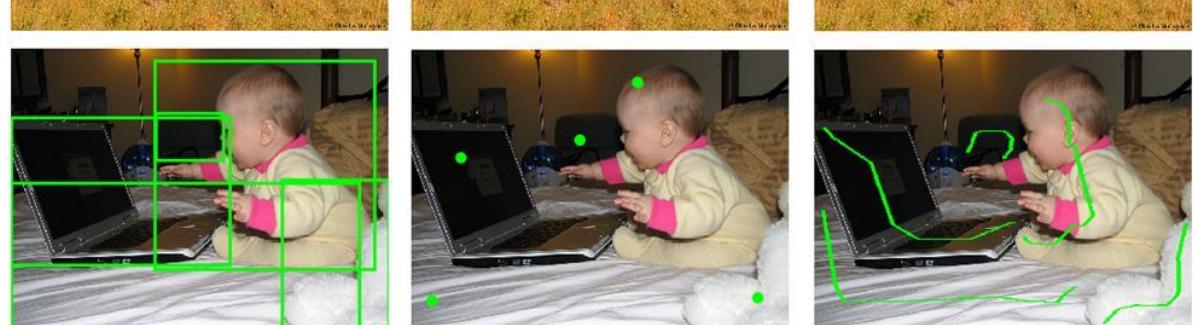
→ Point



→ Scribble

Mask

Skeleton



Boxes

Points

Scribbles

Quantitative Results and Ablation Studies

| Methods | Test-scale | Label | AP | AP-50 |
|------------------|------------|-----------|-------------|-------------|
| PCL [52] | multi | tags | 8.5 | 19.4 |
| C-MIDN [14] | multi | tags | 9.6 | 21.4 |
| WSOD2 [60] | multi | tags | 10.8 | 22.7 |
| Ours | multi | tags | 11.4 | 24.3 |
| Ours | single | tags | 10.8 | 23.1 |
| Ours | single | points | 12.4 | 27.0 |
| Ours | single | scribbles | 13.7 | 29.8 |
| Fast-RCNN [16] | single | boxes | 18.9 | 38.6 |
| Faster-RCNN [43] | single | boxes | 21.2 | 41.5 |
| Ours | single | boxes | 25.7 | 46.3 |

Single model performance on COCO-80(VGG-16)

| Loss | $\mathcal{L}_C + \mathcal{L}_R$ | $+ \mathcal{L}_{T1}$ | $+ \mathcal{L}_{T2}$ | $+ \mathcal{L}_{T1} + \mathcal{L}_{T2}$ | $+ \mathcal{L}_{T1} + \mathcal{L}_{T2} + \mathcal{L}_{Tags}$ |
|-------|---------------------------------|----------------------|----------------------|---|--|
| AP | 22.6 | 24.8 | 25.1 | 25.5 | 25.7 |
| AP-50 | 42.4 | 44.1 | 45.0 | 46.0 | 46.3 |

Strong teacher helps model

| Train | Methods | backbone | Labels | AP | Extra | Labels | AP | Δ |
|----------|---------|-----------|--------|------|---------|-----------|------|-------------|
| COCO-35 | ours | VGG-16 | tags | 4.9 | COCO-80 | - | 5.3 | 8.2% |
| COCO-115 | ours | VGG-16 | tags | 12.9 | Un-120 | - | 13.6 | 5.4% |
| COCO-35 | ours | ResNet-50 | tags | 9.8 | COCO-80 | - | 10.5 | 7.1% |
| COCO-35 | ours | ResNet-50 | boxes | 29.1 | COCO-80 | tags | 29.4 | 1.0% |
| COCO-35 | ours | ResNet-50 | boxes | 29.1 | COCO-80 | points | 30.1 | 5.5% |
| COCO-35 | ours | ResNet-50 | boxes | 29.1 | COCO-80 | scribbles | 30.9 | 6.2% |
| COCO-115 | ours | ResNet-50 | boxes | 32.7 | Un-120 | - | 33.9 | 3.7% |

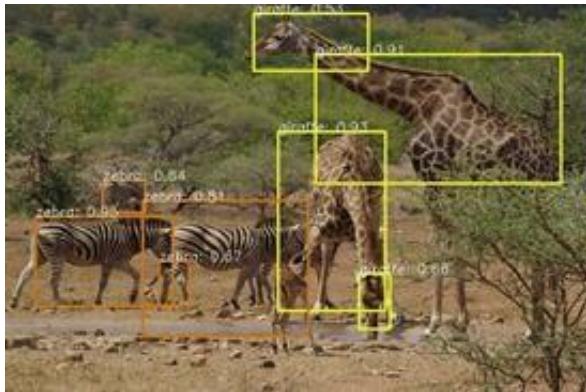
Model improvement with extra data and/or labels

| Methods | Tags | Ref _P | Ref _S | Ref _P + Con _P | Ref _S + Con _S |
|---------|------|------------------|------------------|-------------------------------------|-------------------------------------|
| AP | 10.8 | 11.1 | 11.6 | 12.4 | 13.7 |
| AP-50 | 23.1 | 24.2 | 25.1 | 27.0 | 29.8 |

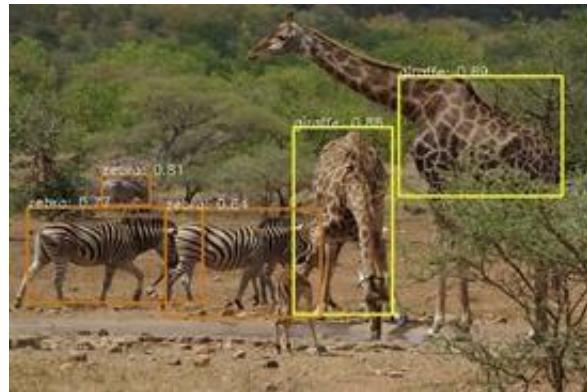
Localization constraints (Con_P/Con_S for points/scribbles) and proposal refinement (Ref_P/Ref_S) helps model

Qualitative Results

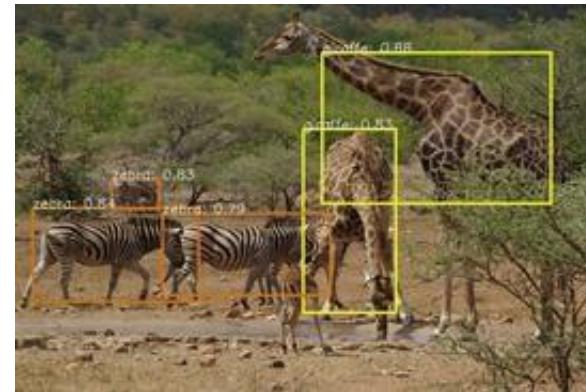
Tags



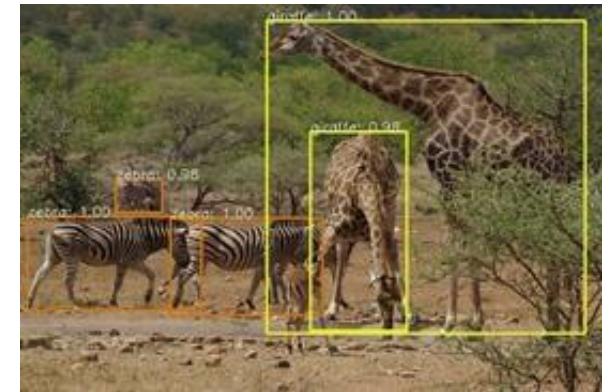
Points



Scribbles



Boxes



person: 0.99
couch: 0.84
person: 0.92
couch: 0.93
person: 0.93
couch: 0.79
person: 0.93
couch: 0.72
book: 0.50
tv: 0.99



A screenshot from a video game showing a room with three people. A bounding box highlights a person's hand holding a remote control with a confidence score of 0.80. Other objects like books, a TV, and a person are also detected with their respective confidence scores.

A screenshot from a video game showing a multi-camera view of a soccer match. The screen is divided into four quadrants by a green frame. In the top-left quadrant, a player in a white shirt is performing a bicycle kick. In the top-right quadrant, a player in an orange shirt is running towards the ball. In the bottom-left quadrant, another player in a white shirt is standing near the ball. In the bottom-right quadrant, a player in an orange shirt is also near the ball. Each player has a bounding box around them with a confidence score: person: 0.99, person: 0.99, person: 0.99, and person: 1.00 respectively. A large yellow bounding box covers the entire ball area. The background shows a stadium with spectators and a scoreboard.

Budget-Aware Omni-Supervised Detection

Approx. per-img budget on COCO:

- Tags: 80s, points: 88.7s, scribbles: 160.4s, boxes: 346s

Given a fixed annotation budget (time), the most common strategy:

- **STRONG**: annotate using only boxes

UFO² shows the possibility of a more efficient new policy:

- N%B: use N% budget for boxes and (1-N%)for points

| Policy | Labels | # Labeled Images | AP |
|---------------|-------------------|-----------------------------|------------------------------------|
| MOST | T | 10000 | 3.00 ± 0.57 |
| STRONG | B + U | 2312 + 7688 | 13.97 ± 0.98 |
| EQUAL | T + P + S + B + U | $2500 + 2255 + 1250 + 5417$ | 5.87 ± 0.70 |
| 80%B | P + B + U | 1804 + 1850 + 6346 | 14.11 ± 1.01 |

Wetectron: A Partially Supervised Learning Platform

The screenshot shows the GitHub repository page for 'NVlabs/wetectron'. The repository has 42 stars, 87 forks, and 5 open issues. It features a README.md file in Python 3.7. The repository description is: "Weakly-supervised object detection platform." It includes tags for 'object-detection', 'weakly-supervised-learning', 'computer-vision', and 'deep-learning'. The 'About' section states: "Weakly-supervised object detection platform." The 'Releases' section indicates "No releases published".

NVlabs/wetectron: Weakly-supervised learning platform

Search or jump to... Pull requests Issues Marketplace Explore

Unwatch 42 Unstar 87 Fork 5

Code Issues Pull requests Actions Projects Wiki Security Insights

master 1 branch 0 tags Go to file Add file Code

jason718 Create README.md 1434905 on Apr 16 1 commits

README.md Create README.md 4 months ago

README.md python 3.7

Wetectron

Wetectron is a software system that implements state-of-the-art weakly-supervised object detection algorithms.

[Project](#) | [Paper](#)

Instance-aware, Context-focused, and Memory-efficient Weakly Supervised Object Detection.
Zhongzheng Ren, Zhiding Yu, Xiaodong Yang, Ming-Yu Liu, Yong Jae Lee, Alexander G. Schwing, Jan Kautz.
In CVPR 2020.

More details coming soon...

Citation

Code Repository

<https://github.com/NVlabs/wetectron>

Conclusions and Future Work



Conclusions

- Eliminating uncertainty is the key
- The diversity of supervisions create challenges but also brings abundant implicit constraints

Future Works

- Multi-modal/multi-task platform vs. single algorithms
- Post-deployment continuous interactive learning
- Inspiration from human baby learning



A network graph is visible on the left side of the slide, consisting of numerous small, semi-transparent green and white circular nodes connected by thin gray lines.

Thank You!

